



Auto-scaling your API

Insights and Tips from the Zalando Team

JBCNConf 2015
Barcelona, Spain

Sean Patrick Floyd - @oldJavaGuy

Luis Mineiro - @voidmaze

ONE of EUROPE'S LARGEST ONLINE FASHION RETAILERS

15 countries

3 fulfilment centers

15+ million active customers

2.2+ billion € revenue 2014

130+ million visits per month

8.000+ employees

Tech hubs in Berlin, Dublin, Dortmund
and Helsinki

Visit us: tech.zalando.com



Our Scale



**2 datacenters
thousands of production instances
serving 15 countries**

Zalando stack





[MUJER](#)[HOMBRE](#)[NIÑO](#)

Iniciar sesión ▾

Favoritos

Cesta

[News & Style](#) [Ropa](#) [Zapatos](#) [Deporte](#) [Bolsos y Complementos](#) [Premium](#) [Marcas](#) [Rebajas](#)

Busca aquí: zapatos, ropa, marcas...



ELIGE UNA SECCIÓN



MUJER



HOMBRE



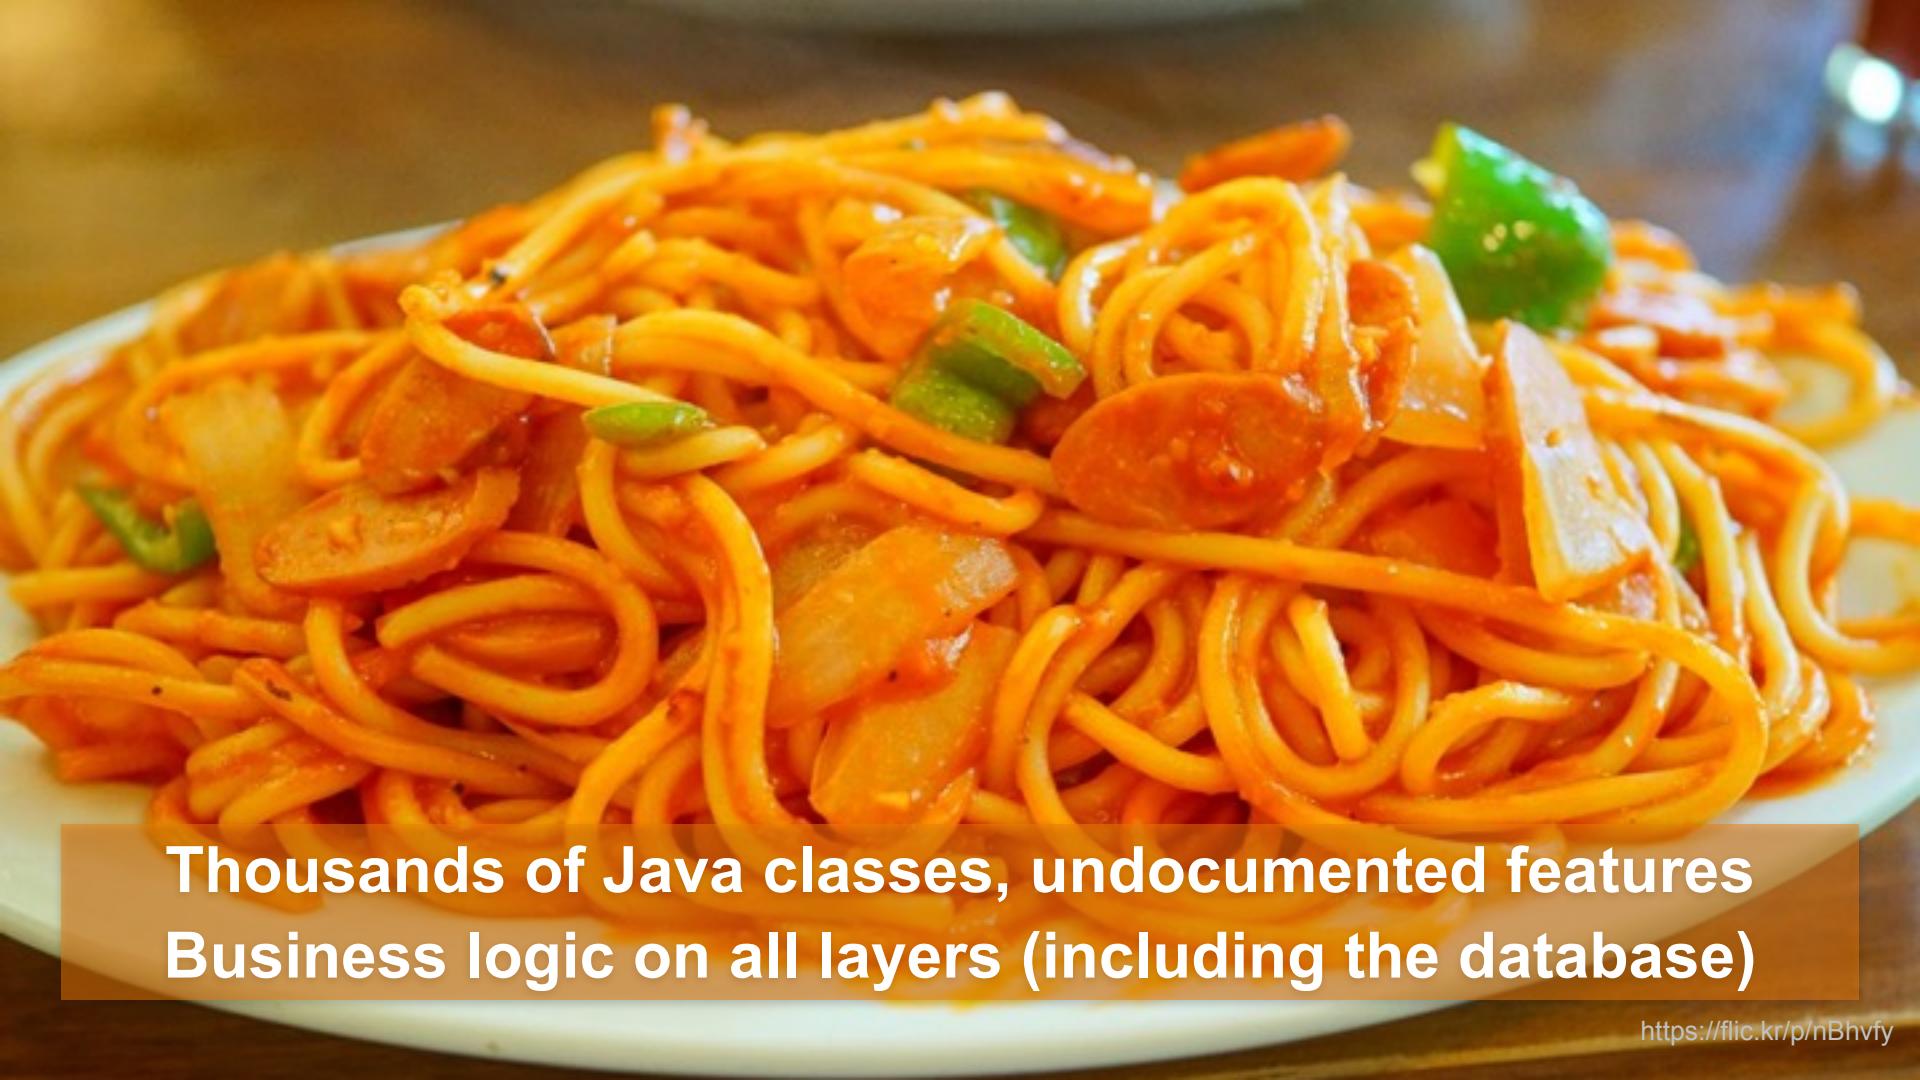
NIÑO



The Shop Monolith

We call it “Jimmy”





**Thousands of Java classes, undocumented features
Business logic on all layers (including the database)**

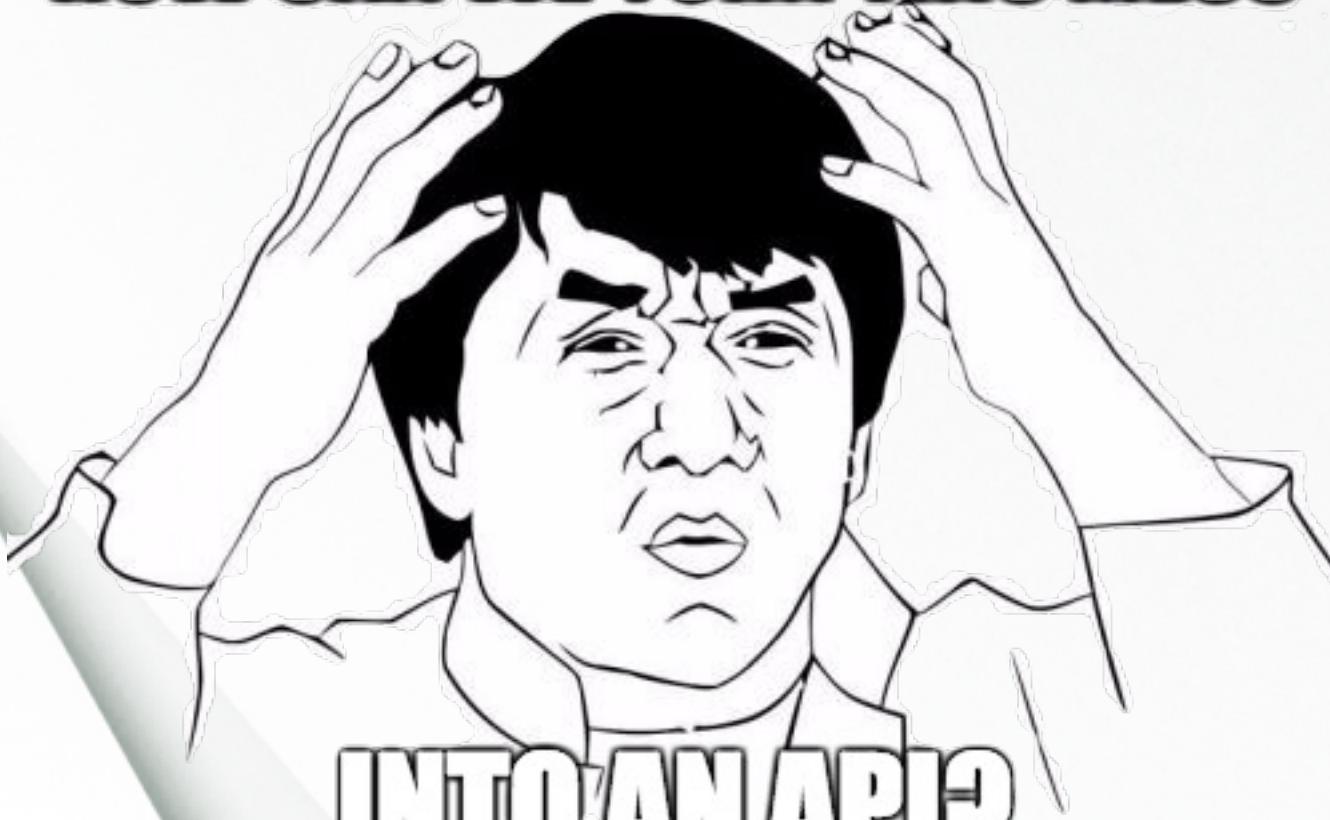


It's 2013...

Zalando wants to
sponsor hackathons...

We need a REST API!

HOW CAN WE TURN THIS MESS



INTO AN API?

Hackathon API

First Step

Let's create some REST-(ish) Spring controllers inside our **monolithic web application**

Deploy a couple of Jimmy instances just to serve requests for this “API”

Pros

The infrastructure is already there!

A few days of coding and we're all set

Cons

This “API” cannot be deployed independently of Jimmy.

Jimmy is infested with business logic everywhere.

New Requirements

WE HAVE A HACKATHON API?



LET'S USE IT FOR SOME OTHER STUFF AS WELL

New Requirements

We needed this new API for our existing frontends, also (very high traffic).

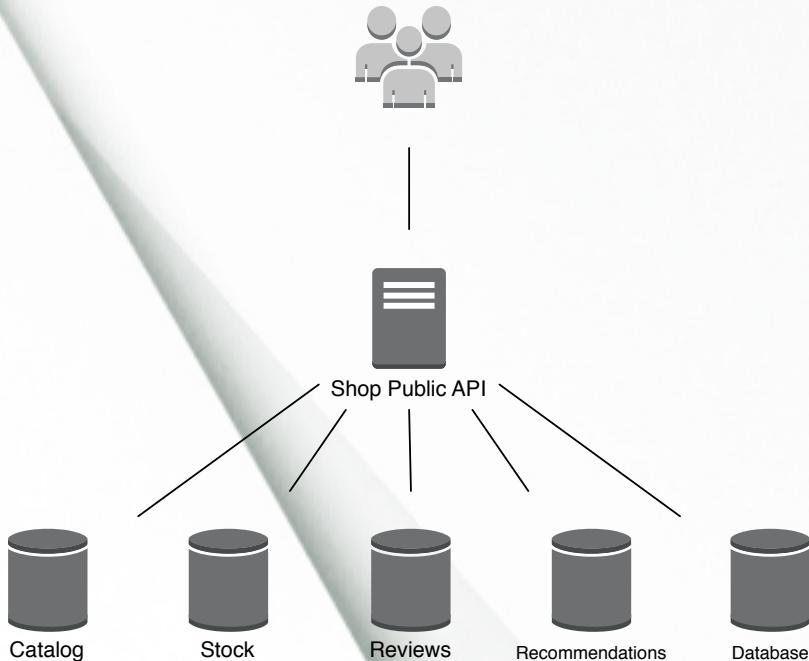
Some third-party apps also wanted to use this API as their backend.

Shop Public API

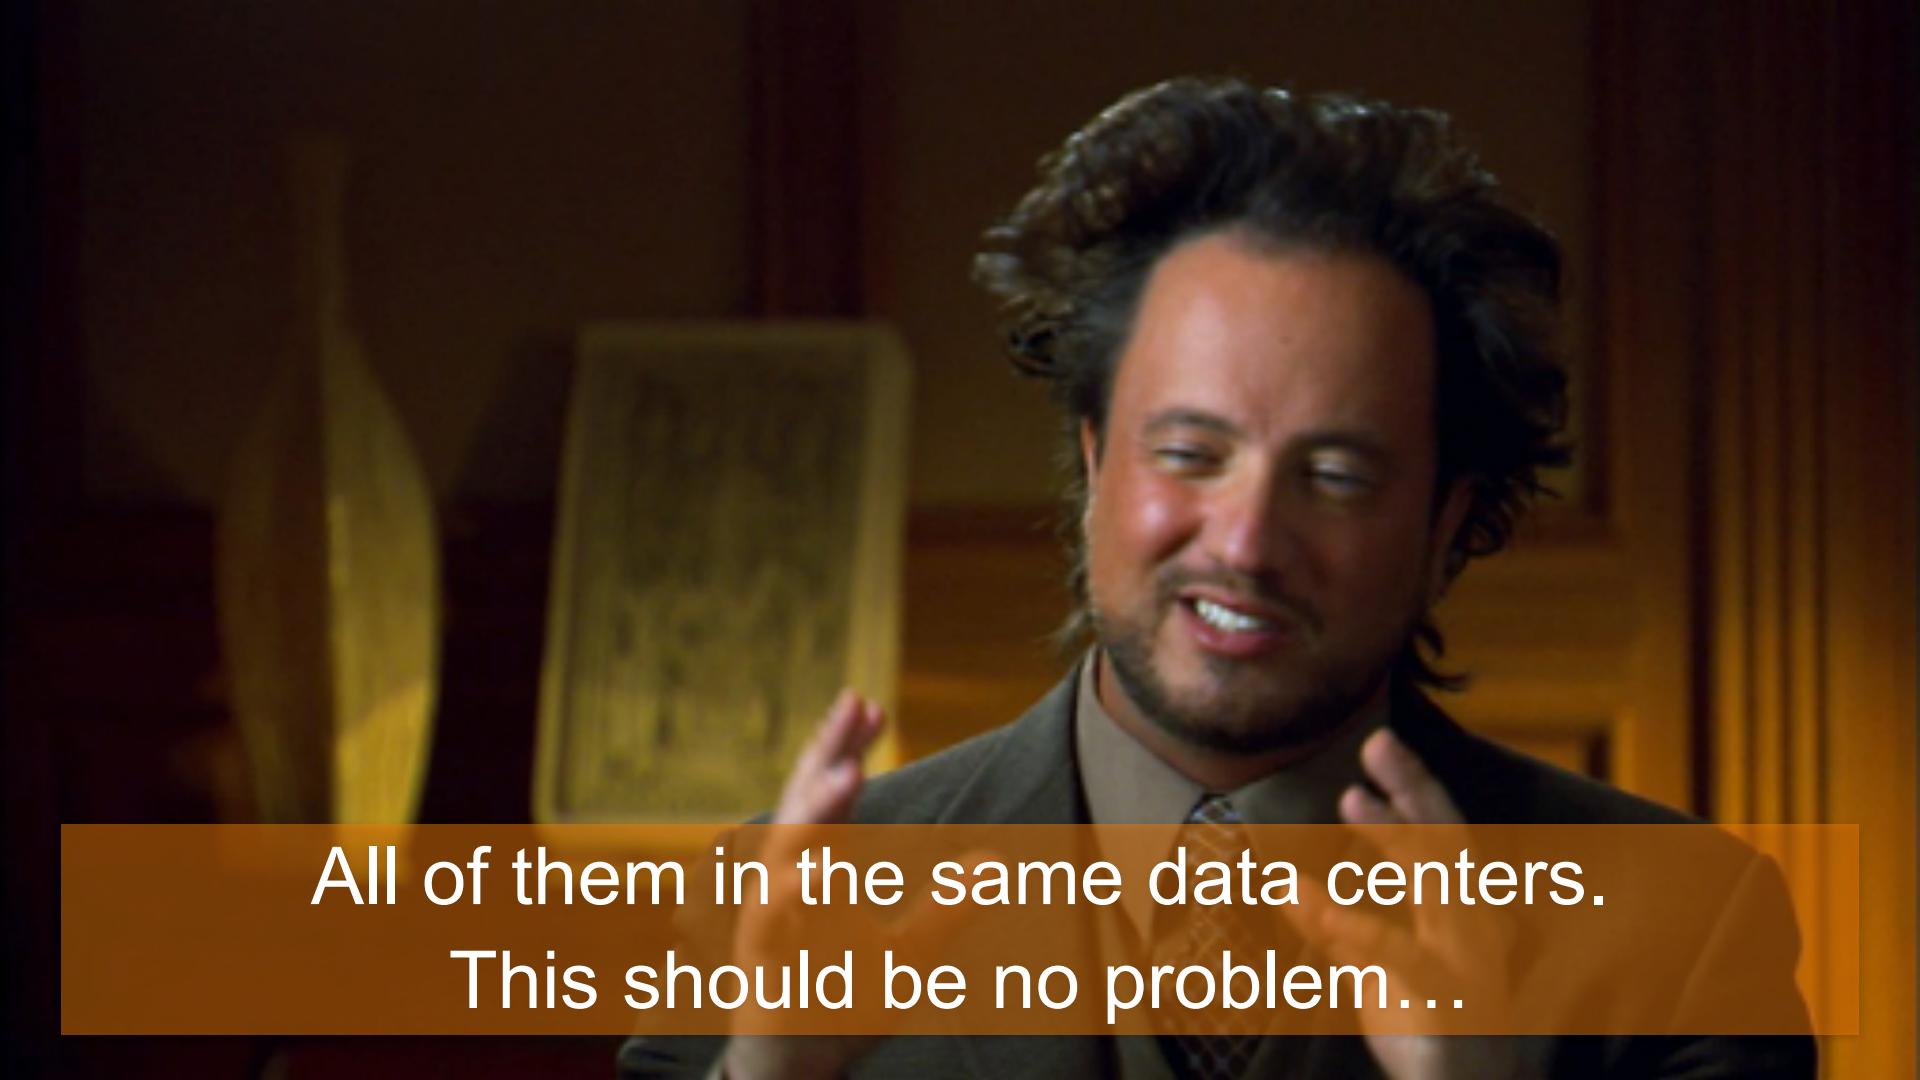
We decided to build a “real” API as a separate standalone project.

Of course, there were some dependencies...

Data Sources



1. Catalog - **SOLr**
2. Stock - **Memcached**
3. Reviews - **REST(-ish) API**
4. Recommendations - **RPC**
5. Database - **PostgreSQL**

A close-up photograph of a man with dark, curly hair and a well-groomed beard. He is wearing a dark suit jacket over a light-colored shirt and a patterned tie. He is gesturing with both hands raised near his shoulders while speaking. The background is a warm-toned wooden interior.

All of them in the same data centers.
This should be no problem...

Tech Company

Paradigm Shift

Until 2014

We're a Fashion Company that has a lot
of tech knowledge

2015

Suddenly we're a Tech Company, providing
“Fashion as a Service”



We established
5 principles

API FIRST

API First

Document and peer review your API in a format like Swagger before writing a single line of code.

Ideally, either generate either your server interfaces or your test data (or both) from the Swagger spec

REST

REST

Manipulate resources, don't call methods

Expose nouns, not verbs

Use HTTP verbs
GET, PUT, POST, DELETE, PATCH

SAAS

Identity Management

Our backend services don't expose APIs yet

Company-wide IAM strategy is not ready yet

Can only expose read-only features for now

MICRO- SERVICES

Microservices

We already have a Service-Oriented
Architecture

It was mostly SOAP, though...

And definitely not micro

CLOUD

Road to AWS

Some challenges ahead...

1. Backend services not available yet
2. SOLr also not available
3. We can't just move the databases there

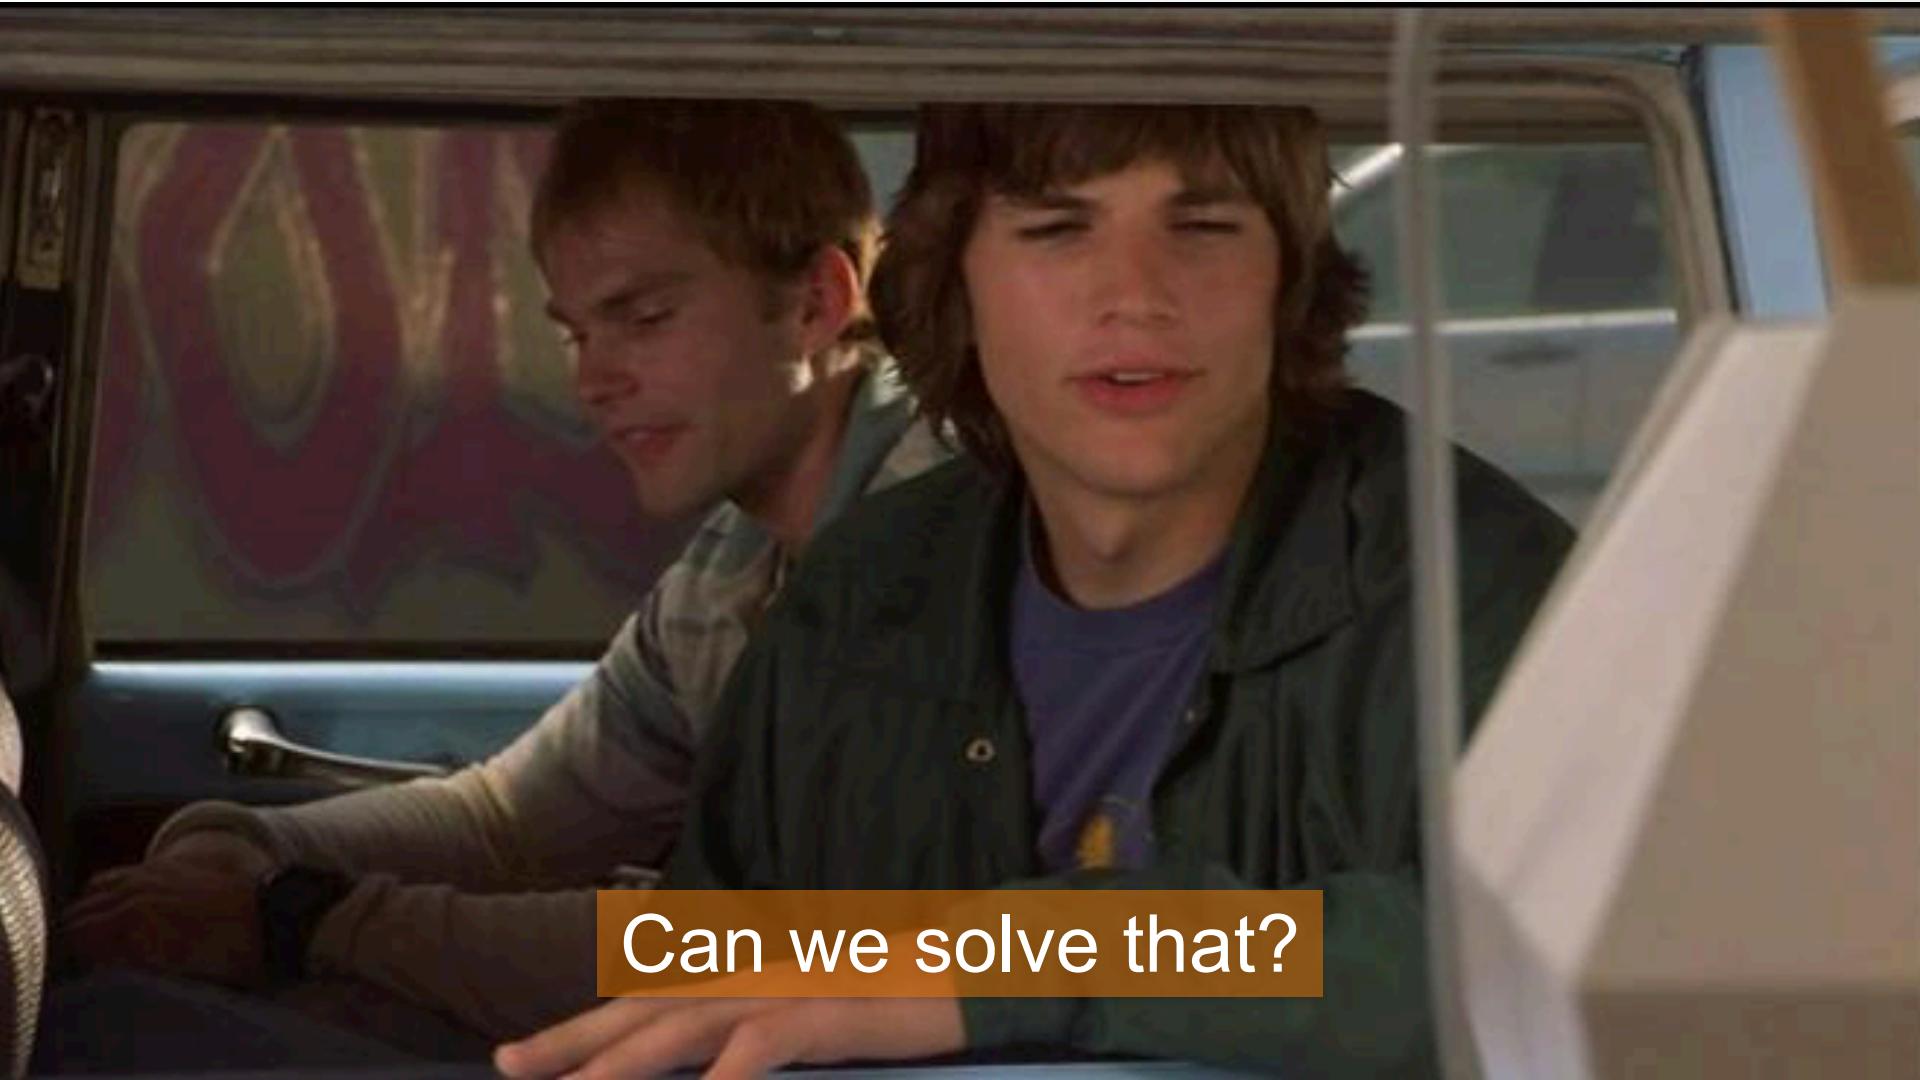
How we did it

Challenges

Catalog and Stock datasources are latency critical



and they're owned by different teams!

A scene from the movie "The Social Network". Two young men, one with short brown hair and another with long dark hair, are sitting in the back seat of a car. They are looking at a laptop screen which is partially visible on the right side of the frame. The man with long hair has his eyes closed and is speaking. The man with short hair is looking down at the laptop. The background shows the interior of the car and some blurred trees outside.

Can we solve that?



Step 1 - move critical datasources to AWS

“Move” Data Sources to AWS

We can't just **move** them!

Jimmy also needs them in the data centers,
you insensitive clod!

Ok... we just **replicate** them.



STEAMED RICE

SOUUPS

WONTON SOUP

EGG FLOWER SOUP

HOT AND SOUR SOUP

BEEF NOODLE SOUP

CHINESE
FOOOOOD

and then...?

Replicating Data Sources

SOLr has its own replication mechanism
over HTTP

memcached should be easy ...

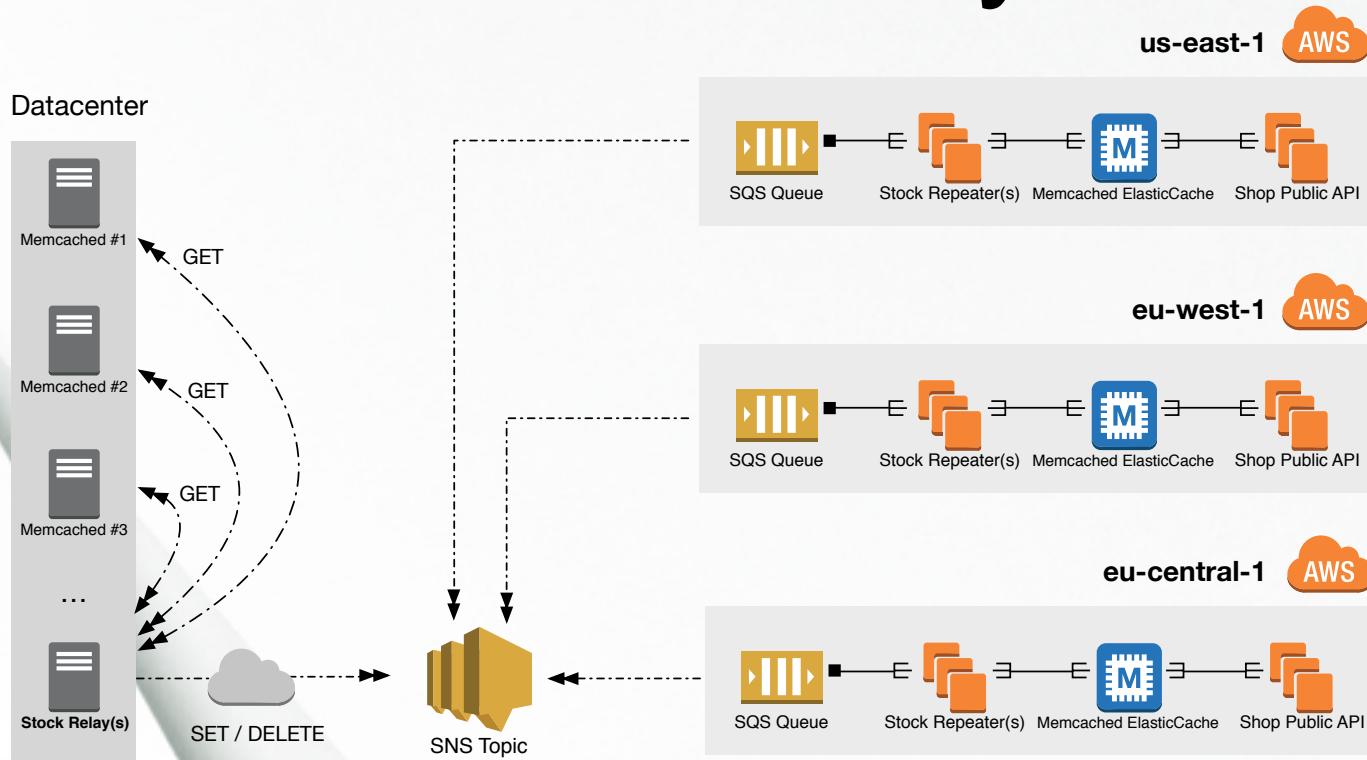
A photograph of five women in black suits standing outdoors. They are all looking directly at the camera with serious expressions. The background consists of green bushes and a small white decorative structure.

You wish!

Step 2 - Build memcached replication



Stock Relay





STEAMED RICE

SOUUPS

WONTON SOUP

EGG FLOWER SOUP

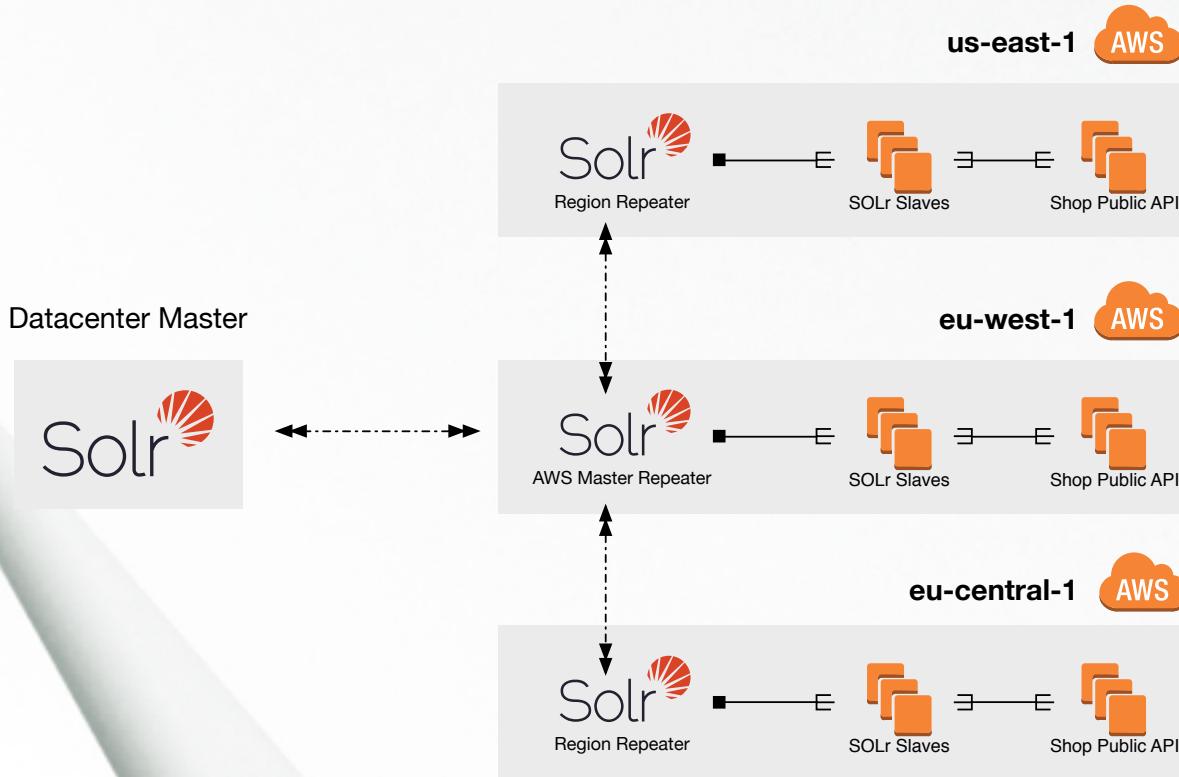
HOT AND SOUR SOUP

BEEF NOODLE SOUP

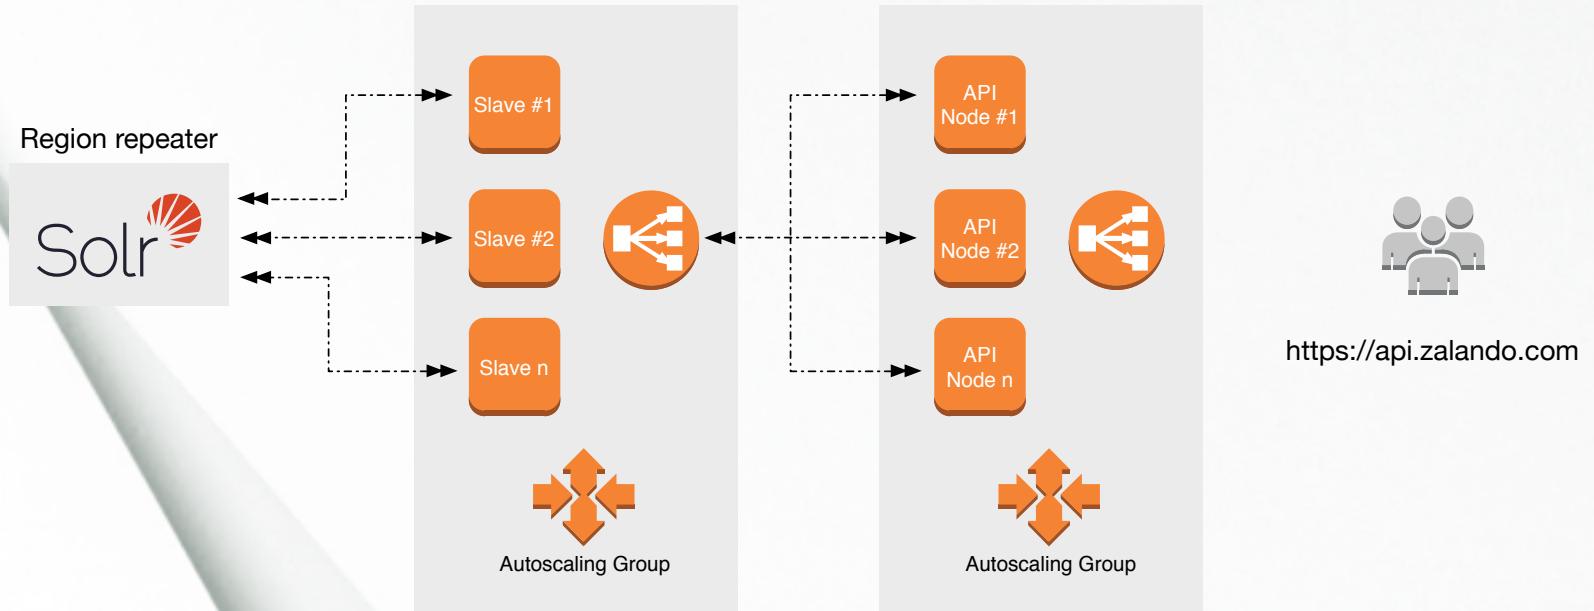
CHINESE
FOOOOOD

and then...?

AWS SOLr Repeater

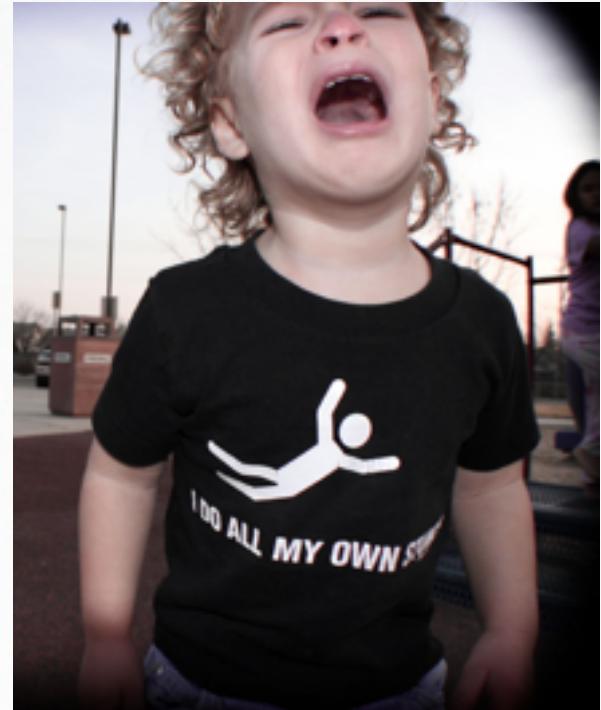
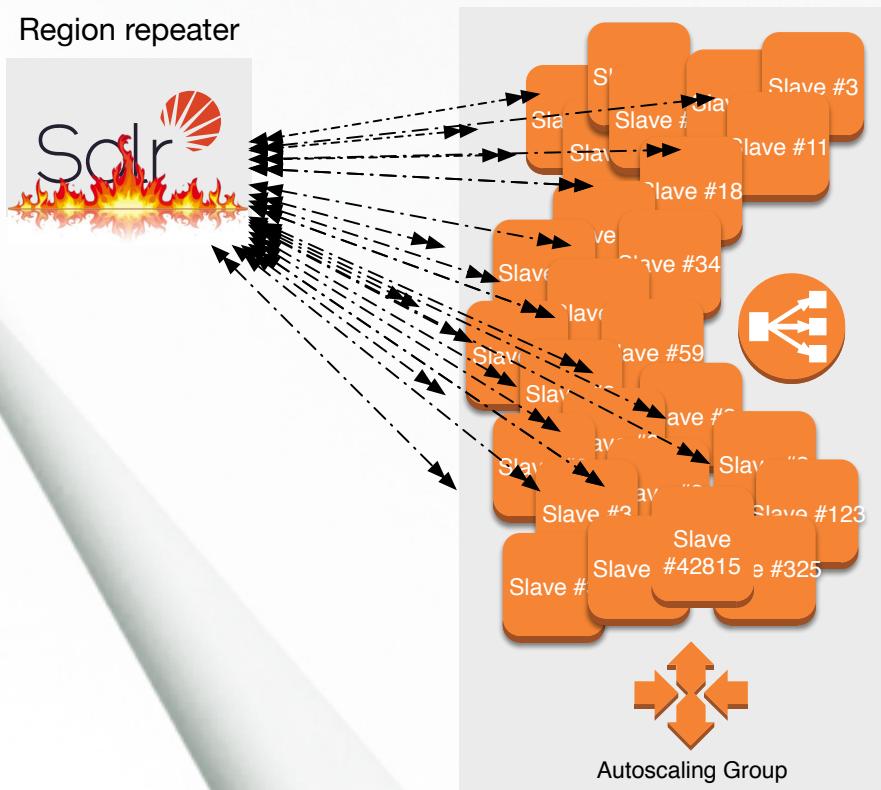


Autoscaling SOLr and API



Scaling Limitations

Region repeater

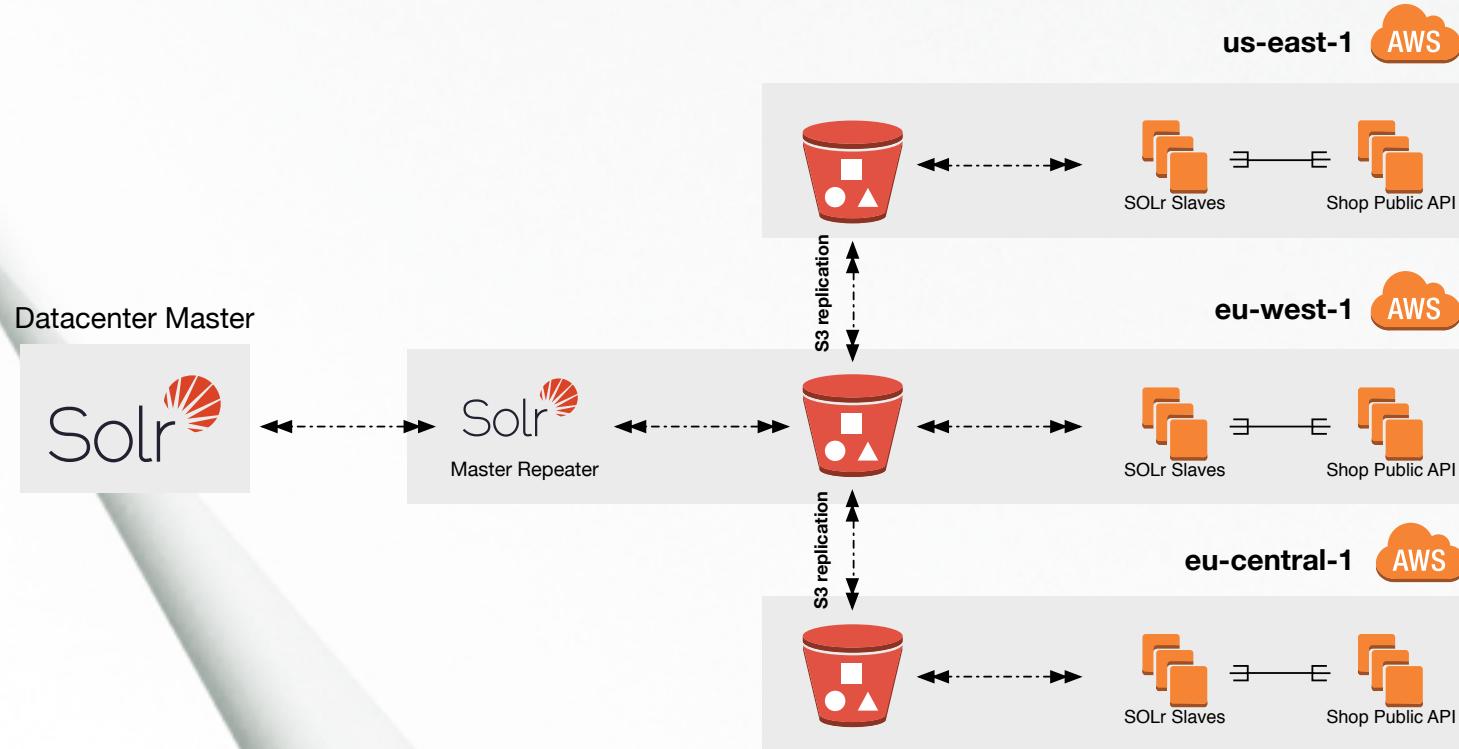


Edited from <https://flic.kr/p/zuBXM>

A man with brown hair, wearing a green camouflage long-sleeved shirt, is leaning over the front of a light-colored car. He is looking down at something in his hands, possibly a tool or a part. The background is dark and out of focus.

Y U no scale!

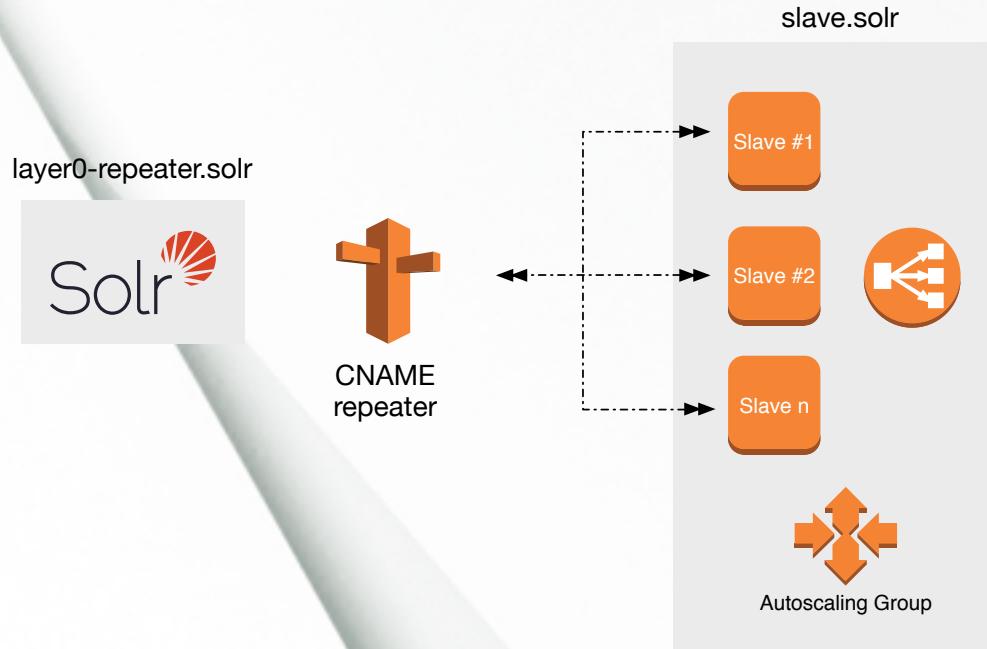
S3 Bucket for Replication



Spec

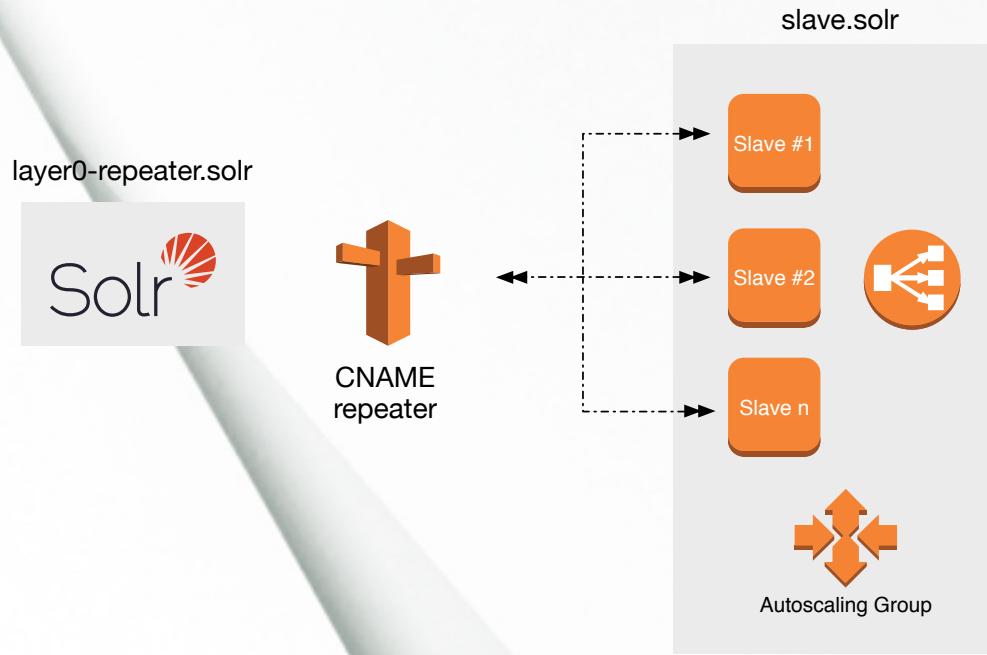
I see...

“Onion layers” for Replication



- Start with a single repeater layer - **layer0.solr**
- Setup a Route53 CNAME repeater that links to it - **repeater.solr**
- Entire slave fleet in the ASG is always configured to replicate from **repeater.solr**

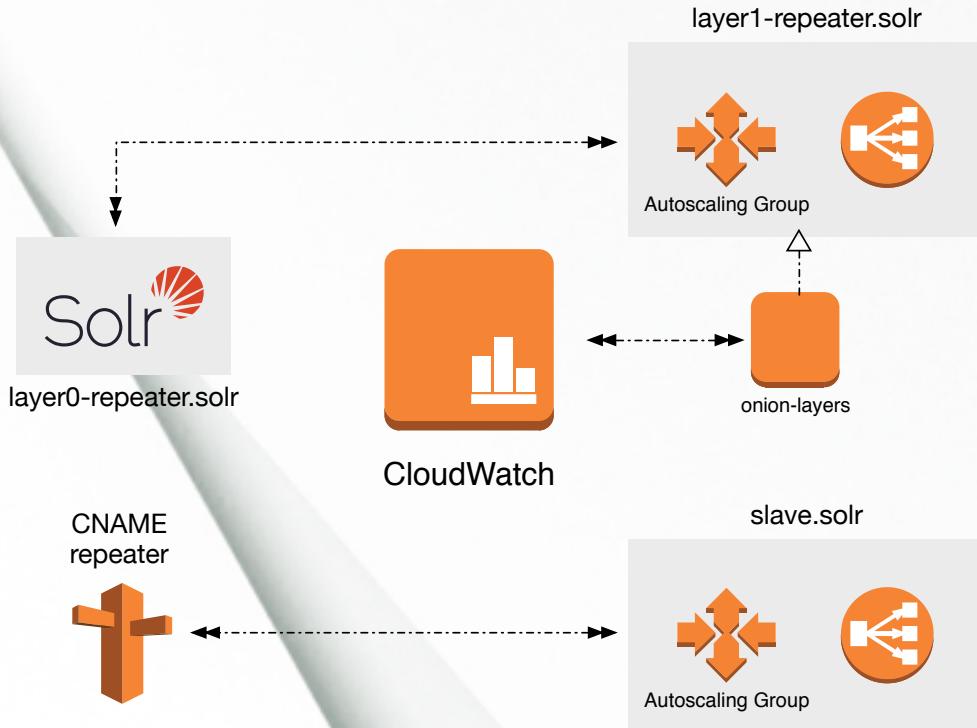
“Onion layers” for Replication



CloudWatch

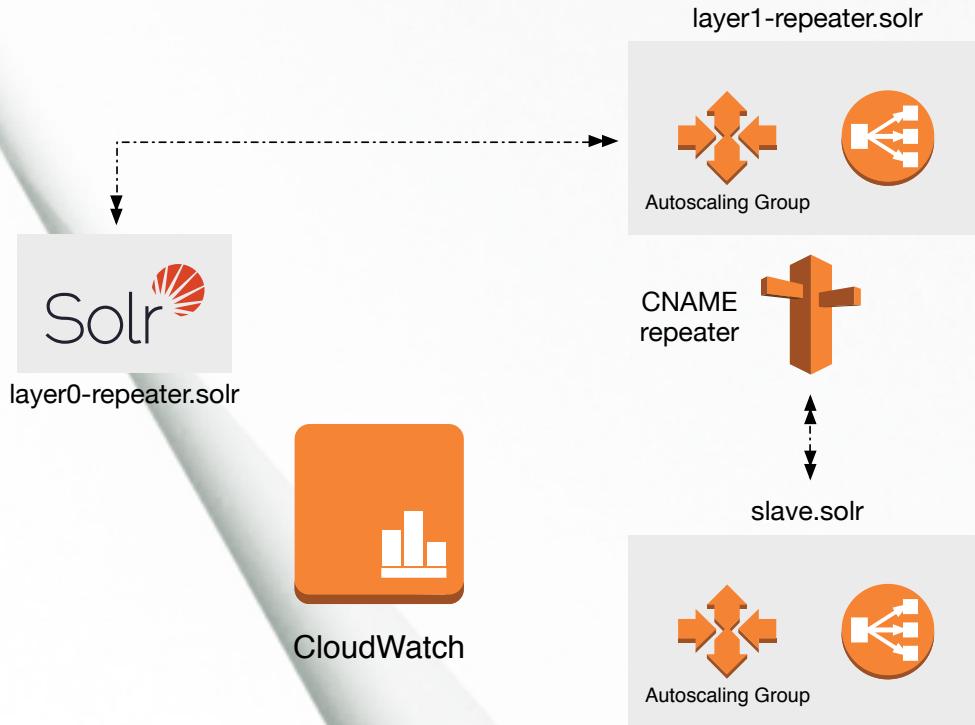
- Setup CloudWatch alarms for relevant metrics in the repeater layer. Give them some slack space.
- Configure it to send notifications to the **onion-layers-topic**.
- Bring up your instance of **onion-layers**.

“Onion layers” for Replication



- **onion-layers** creates a new ASG layer. Calculates ***currentLayer* = *currentLayer* + 1**
- The new layer starts replicating from **layer<currentLayer-1>-repeater**
- Adds a new Route53 recordset **layer<currentLayer>-repeater**.

“Onion layers” for Replication



- After replication, the CNAME **repeater** is updated to link to **layer<currentLayer>-repeater**.
- **onion-layers** acts on alarms for the connection between **current layer** and **previous layer**.
- You can still have your own AS alarms for the connection between **slaves** and **current layer**.

A scene from the TV show 'Two and a Half Men' featuring Charlie Sheen and Ashton Kutcher. Charlie Sheen, on the right, has blonde hair and is wearing a striped shirt. Ashton Kutcher, on the left, has brown hair and is wearing a dark green jacket over a blue t-shirt. They are both looking towards the camera with slight smiles. The background is a dark room with a green sign that partially reads 'ERRA'.

Yes. These tools will be open-sourced soon

Thank you for listening

Check out our blog <https://tech.zalando.com>

Our many open source products <https://github.com/zalando>

The STUPS stack <https://stups.io>

Got more questions? You can reach us on twitter

@ZalandoTech

We're hiring!

Special thanks to Jessie Dude.

No Continuum Transfunctioners were harmed during the production of these slides.