



ZMON - Monitoring Our Platform

DevOps Meetup Dublin | September 3, 2015 | jan.mussler@zalando.de | [@JanMussler](https://twitter.com/JanMussler)

ONE of EUROPE'S LARGEST ONLINE FASHION RETAILERS

15 countries

3 fulfillment centers

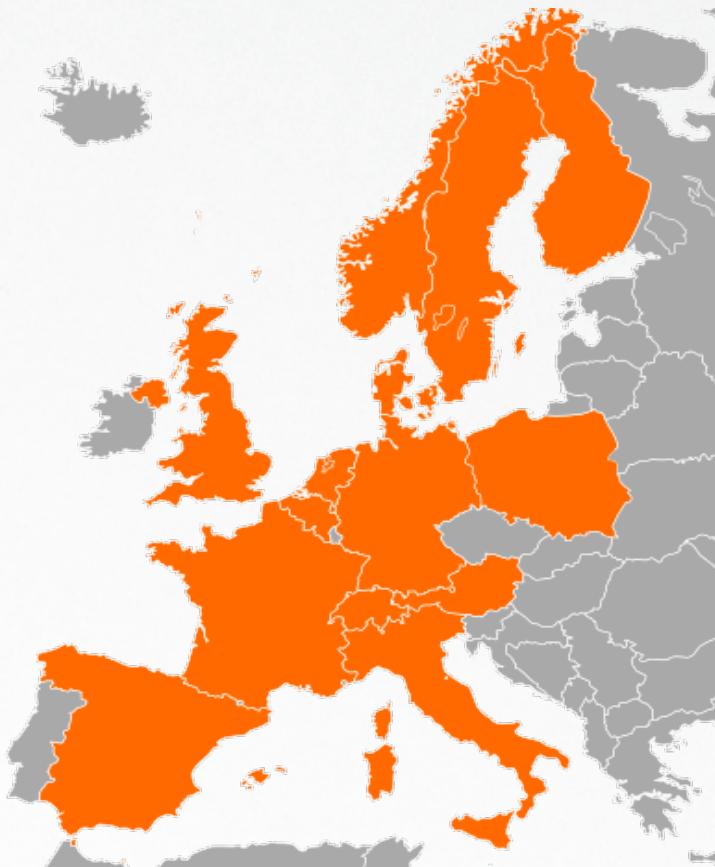
16+ million active customers

2.2+ billion € revenue 2014

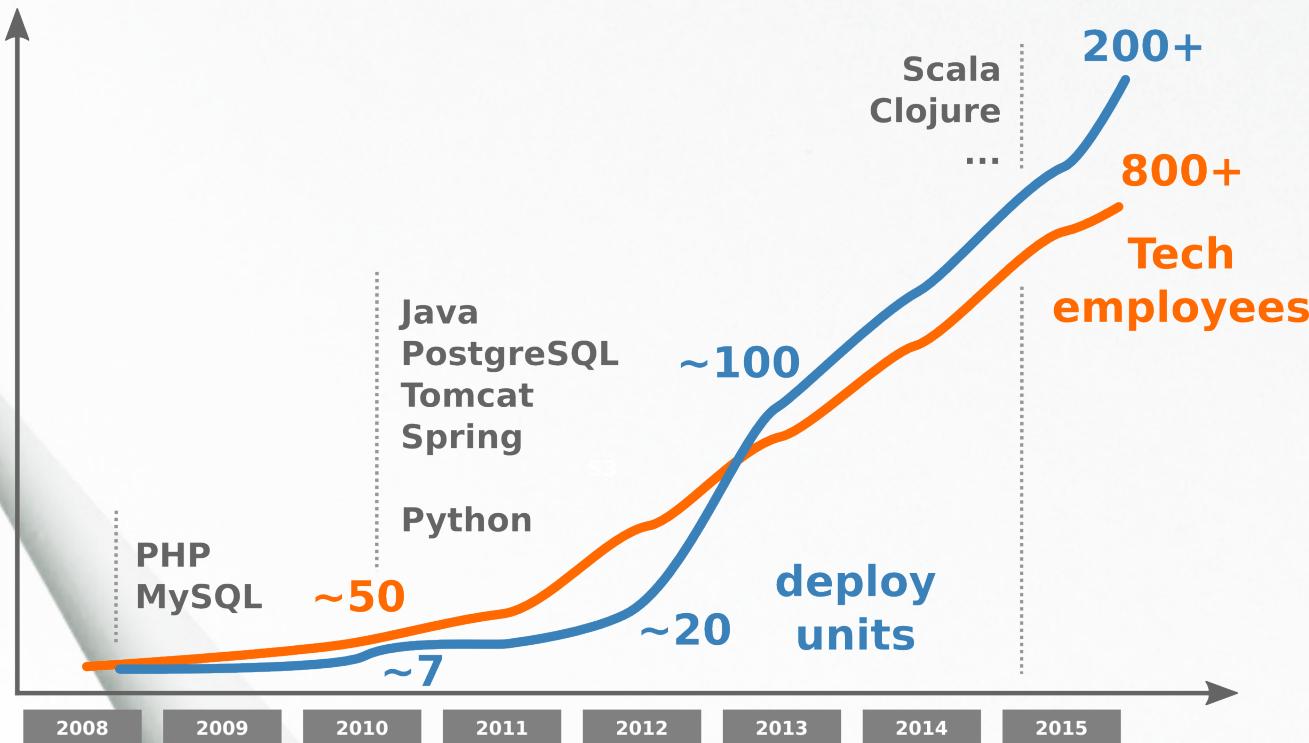
130+ million visits per month

8.000+ employees

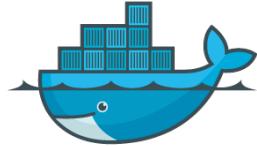
Visit us: tech.zalando.com



Zalando's Technology History



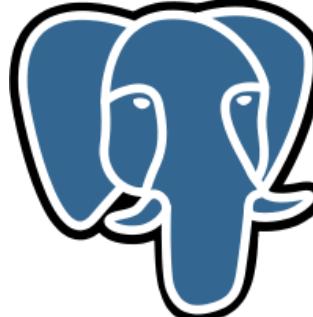
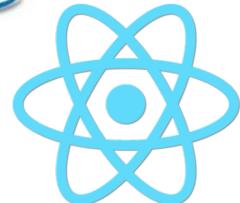
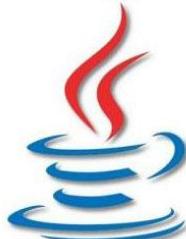
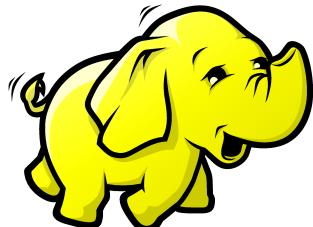
(Some!) Technologies We Use



docker



ANGULARJS



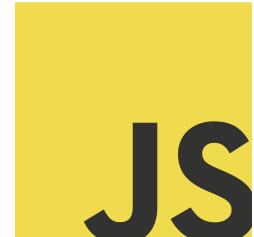
zalando



cassandra



python



Monitoring Situation Until Late 2013

ICINGA plus custom frontend (ZMON 1)

Did not scale with growth:

- Our UI became too slow
- Number of systems to check too many
- Number of teams that wanted checks grew
- Every request had to go through single team

Goals of new ZMON development

Improve performance and throughput

Autonomy for individual teams

Flexibility and extendability

Integration into tooling (CMDB, DeployCtl ...)

The basic terminology ...

Entity:

Anything you may want to monitor

Can be used as a "dimension"

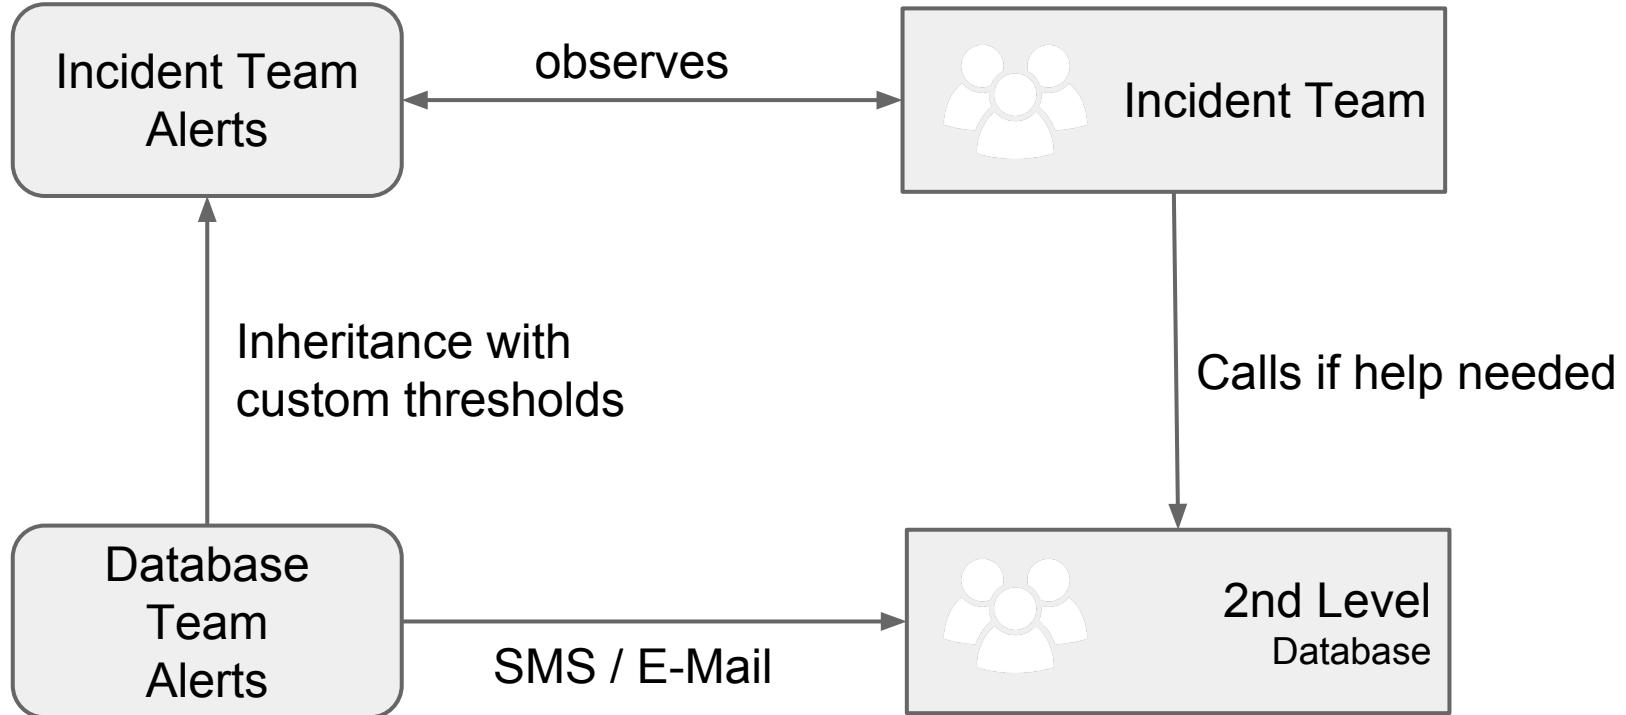
Checks:

Runnable Python snippet fetching data

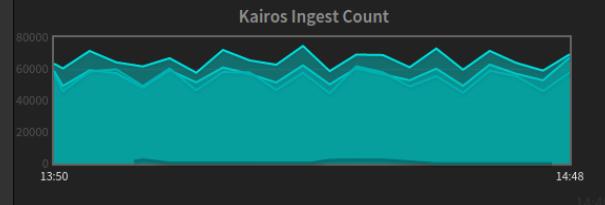
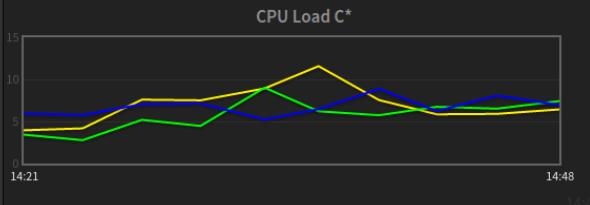
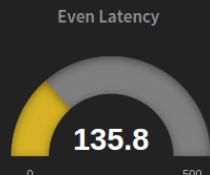
Alert on Check:

Python expression yielding true or false

Zalando Tech - 24x7 team setup



Customizable ZMON dashboards



Filter alerts

0 Hide Widgets

STUPS Application Errors

a-twintip-crawler[701] (52.79)



STUPS Infinity

STUPS ELB Request 500

elb-pierone-b63[701] ({"requests_per_sec":0.3833333333333336,"http5xx_errors_per_sec":0.01666666666666666,"http4xx_error...")



STUPS Infinity

ZMON Snmp Queue Too Large

monitor03 (1163)



STUPS 28s

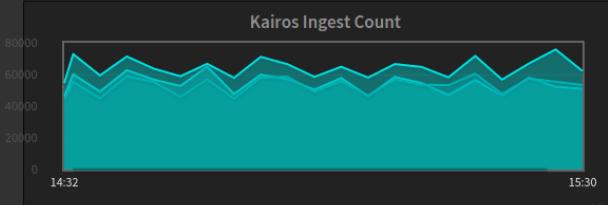
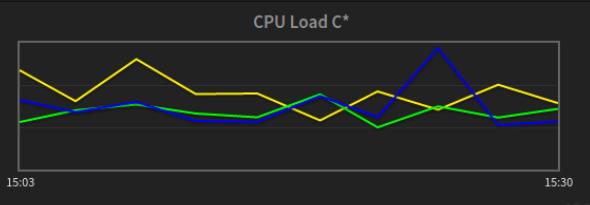
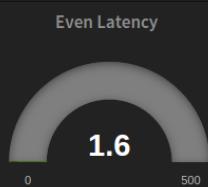
EventLog FinalErrors: 1

itr-elsn03:5827 (1)



STUPS 13m

Customizable ZMON dashboards



Filter alerts

0 Hide Widgets



STUPS Application Errors

STUPS ELB Request 500 (2)

EventLog FinalErrors: 1, 2 (4)

⚠ Cassandra Write Timeouts (3)

⚠ {c} Cassandra Nodes down (3)

⚠ Cassandra Read Timeouts (3)

⚠ Low Free Memory: {freemem} MiB

ZMON AWS Worker outdated (7)

ZMON AWS Agent outdated (6)

ZMON AWS Scheduler outdated (2)

Check definition import failed

ZMON Redis

Low Free Diskspace

Platform/Software High Load: 30.79, 93.18 (3)

Customizable ZMON dashboards

Dashboard name

Jan's EventLog

Default view

Full

Edit mode

PRIVATE (only you can edit the dashboard)

Filter Tags

No Tags selected

Widgets Configuration



```
1 [  
2 {  
3   "type": "gauge",  
4   "title": "Event Latency",  
5   "checkDefinitionId": 2027,  
6   "entityId": "aws-ac[eu-west-1]",  
7   "jsonPath": ".latency_ms",  
8   "options": {  
9     "max": 500  
10  }  
11 },  
12 {  
13   "type": "chart",  
14   "title": "CPU Load C*",  
15   "options": {  
16     "yaxis": {  
17       "min": 0  
18     },  
19     "metrics": [  
20       {  
21         "label": "CPU Load C*",  
22         "series": "cpu_load_c_star",  
23         "value": 100  
24       }  
25     ]  
26   }  
27 }]
```

Alert Teams

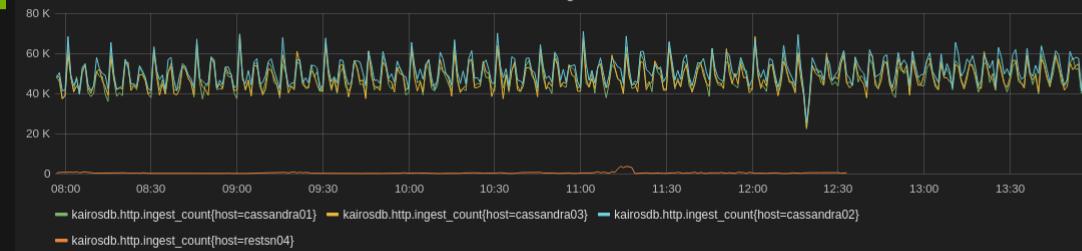
```
1 [  
2   "Platform/Software*",  
3   "STUPS*"  
4 ]
```

Display historic data using Grafana

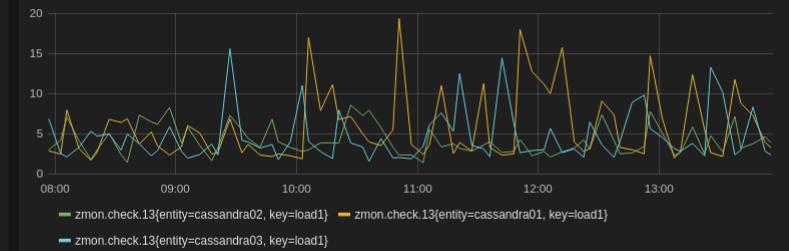
Jan's Dashboard

Zoom Out 6 hours ago to a few seconds ago refreshed every 1m

KairosDB Ingest Count



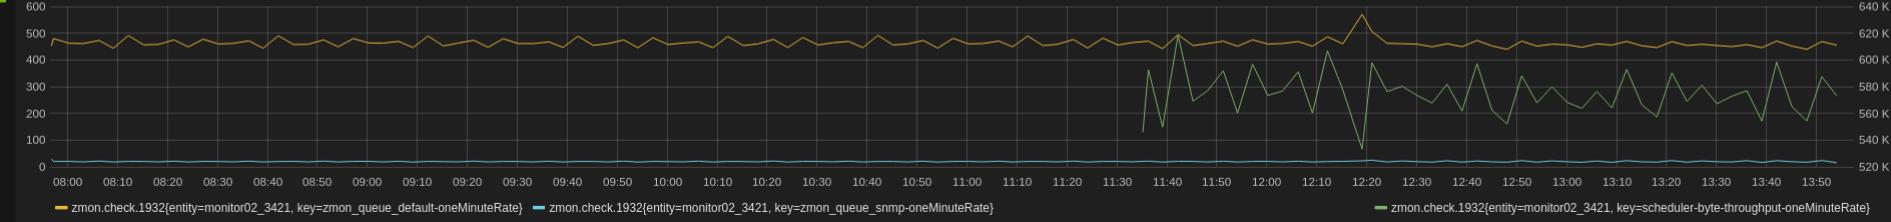
Cassandra Load



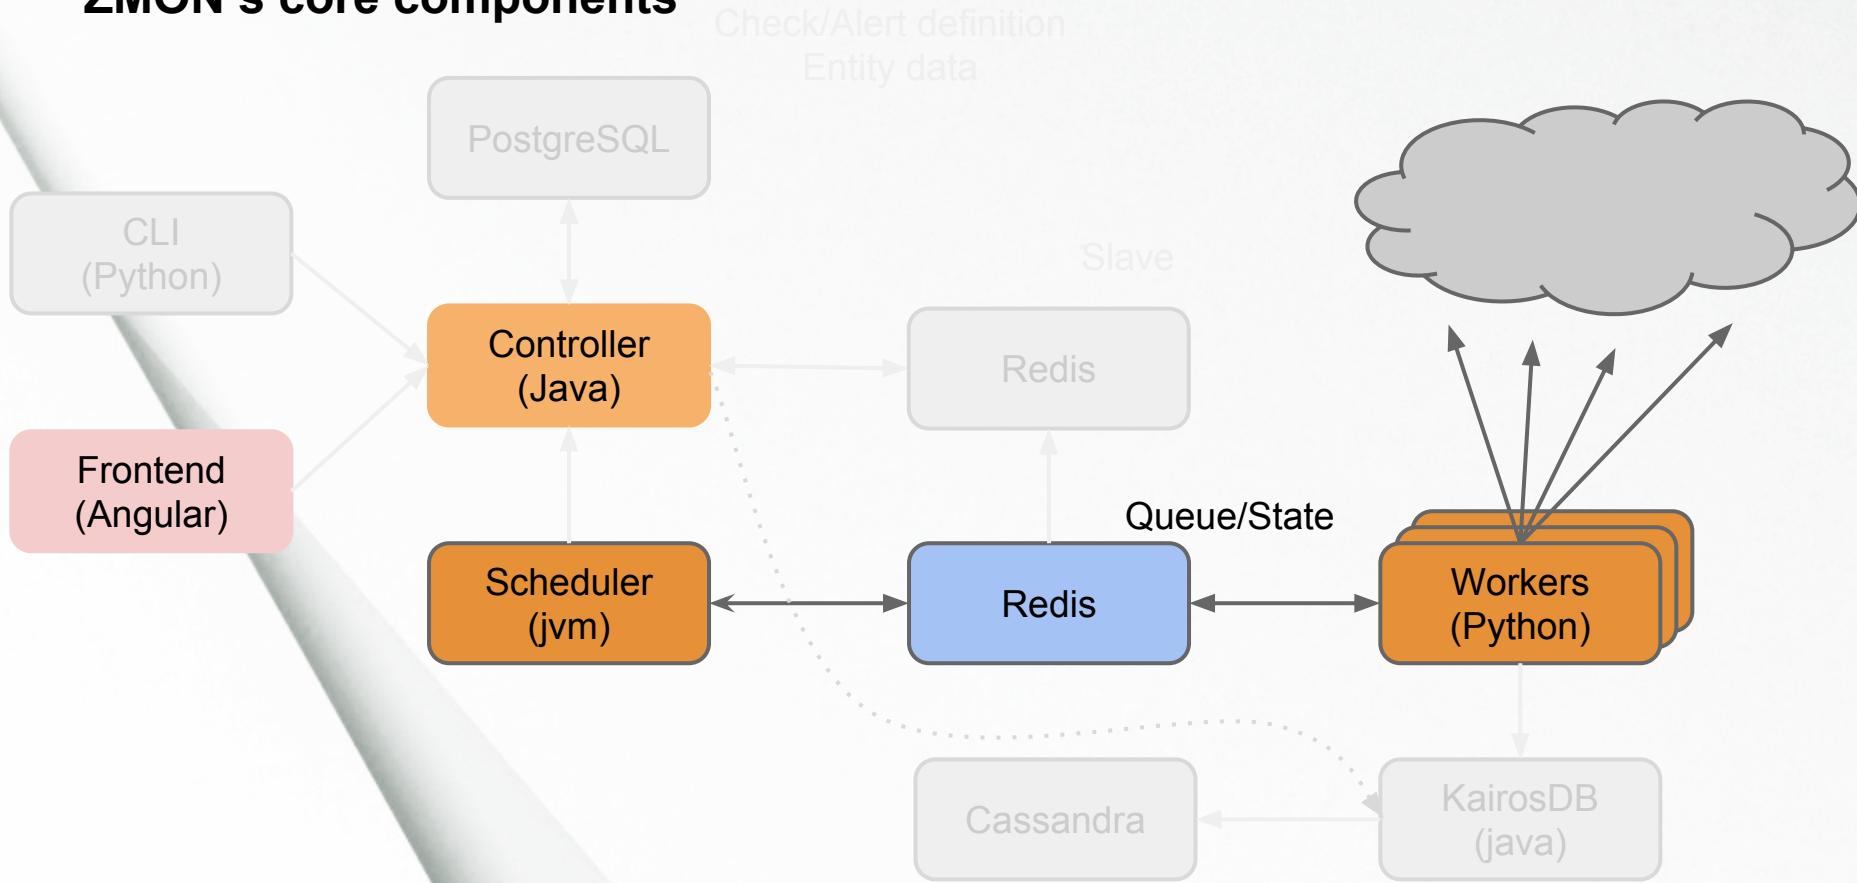
Python Scheduled Checks



Scheduler NG Check / Data Rate



ZMON's core components



Entities

- hosts, databases, applications, instances ...
- generic key value object
- 4000+ entities in our deployment

```
{  
  "id": "node01:8080",  
  "type": "instance",  
  "host": "node01",  
  "ports": {"8080":8080,"8181":8181},  
  "application_id": "zmon",  
  "application_version": "0.1.0",  
  "dc": "dc1"  
}
```

Entity "node01:8080"

Database Entity

```
{  
  "id": "customer-live-slave",  
  "type": "database",  
  "role": "slave",  
  "environment": "live",  
  "shards": {  
    "customer1": "customer1.db:5432/customer1"  
    "customer2": "customer2.db:5432/customer2"  
    "customer3": "customer3.db:5432/customer3"  
    "customer4": "customer4.db:5432/customer4"  
  }  
}
```

Entity: customer-live-slave

Entity Service

Integrated easy-to-use entity store with REST API

```
>zmon entities push local-postgres.yaml
```

local-postgres.yaml

```
id: localhost:5432
type: postgres
host: localhost
port: 5432
shards:
```

```
local_zmon_db: "localhost:5432/local_zmon_db"
```

Checks

- select subset of entities
- executes Python expression
 - powerful using **eval** with custom context
 - Builtins: HTTP, PostgreSQL, MySQL, Cloudwatch, Redis, SNMP, tcp, SOAP, Scalyr...
- returns "value" object
 - Quickly, every check returned "dicts"

Managing checks

REST API to update / auto-import from SCM

```
zmon check-definitions update select-1-check.yaml
```

```
name: "Select 1"
owning_team: "Team 1"
command: |
    sql().execute("select 1 as a").results()
entities:
- type: postgres
interval: 15
description: "test connection"
```

select-1-check.yaml

Check: Number of events written by queue

Description	Number of events written by queue
Command	<pre>def check(): items = http(url='/ws/dispatcher/queuestatus').json().iteritems() r = {} for (tk, tv) in items: for (k, v) in tv.iteritems(): r[k]=v["written"] r["_use_scheduled_time"] = True return r</pre>
Interval	1m
Entities	environment = live AND project = eventlog-dispatcher AND type = zomcat

 Details

[+ Add New Alert Definition](#)


Filter alert definitions

Actions	Name	Condition	Entities	Team	Responsible Team	Priority	Status
 	Number of Events written	False	(no entity filter)	Platform/Software	Platform/Software	?	✓

Alerts

- Executes using a check's value, bound to single check
- Defines team and responsible team
- Allows inheritance from other alert
- Evaluates Python expression yielding True/False
- No "WARNING" state, no "UNKNOWN" state
- Priorities and tags

ID: 123, Status: OK, Team: Platform/Software

Alert: Cassandra Node Load ➔

[Edit](#) [Clone](#) [Inherit](#) [Delete](#) [History](#) [Trial Run](#) [Evaluate](#) [Report an Incident](#) [Comments \(0\)](#)

Description

Load high on Cassandra Nodes

Condition

value["load5"]>13

Responsible Team Platform/Software

[Details](#)[Alerts/Checks](#)[Downtimes \(0\)](#)[History](#)[Children \(0\)](#)[Alerts \(0\)](#)[In downtime \(0\)](#)**OK (13)**

Filter entities

Actions		Name	Timestamp	Last run	Value	Captures
	<input type="checkbox"/>	cassandra01	2015-06-12 18:02 (1m ago)	66ms	{"load1":2.43,"load15":3.46,"load5":3.29}	{}
	<input type="checkbox"/>	cassandra02	2015-06-12 18:02 (1m ago)	61ms	{"load1":3.17,"load15":3.01,"load5":3.35}	{}
	<input type="checkbox"/>	cassandra03	2015-06-12 18:02 (1m ago)	34ms	{"load1":0.91,"load15":1.85,"load5":1.47}	{}

Alert: {c} Cassandra Node Down

 Edit  Clone  Inherit  Delete  History  Trial Run  Evaluate  Report an Incident
 Comments (0)

Description

At least one Cassandra Node is down in Live Cluster

Analysis

You may use nodetool on multiple hosts in the cluster to verify that nodes are down(D). however, node may appear up(U) when running nodetool on the same machine, but node is still unavailable to cluster.

```
nodetool status
```

```
Datacenter: datacenter1
```

```
=====
```

```
Status=Up/Down
```

```
|/ State=Normal/Leaving/Joining/Moving
```

-- Address	Load	Tokens	Owns (effective)	Host ID	Rack
------------	------	--------	------------------	---------	------

Condition

```
capture(c=value["DownEndpointCount"]) > 0
```

Responsible Team

Platform/Software

 Details**Check Name, ID, History**

[Cassandra Cluster Status](#), ID: 501, [History](#)

Check Team

Platform/Software

Check Command

```
jmx(port=7199).query('org.apache.cassandra.net:type=FailureDetector', 'DownEndpointCount', 'UpEndpointCount').results()
```

Check Interval

1m

Check Entities

host_role_id = 104 AND type = host

Trial Run - Quick feedback and download YAML

Details Run

Name Unauthorized Access to Event Data Service {rate}

Check Command

```
13 if ks[-2]=='snapshot':
14     ep = '.'.join(ks[4:-2])
15 else:
16     if isinstance(v, dict):
17         ep = '.'.join(ks[4:])
18     else:
19         ep = '.'.join(ks[4:-1])
20
21     if not ep in r:
22         r[ep] = {}
23
24     # zmon.response. 200 . GET . EP .
25
26     if ks[3] not in r[ep]:
27         r[ep][ks[3]] = {}
28
29     if ks[2] not in r[ep][ks[3]]:
30         r[ep][ks[3]][ks[2]] = {}
31
32
```

Valid Python expression used to query the instances specified in entities filter.

Alert Condition

```
1 def alert():
2     rate = 0
3     for k,v in value.iteritems():
4         for hm, vs in v.iteritems():
5             ks = vs.keys()
6             if '401' in ks:
7                 rate += vs['401']['mRate']
8
8     return capture(rate=rate) > 0.1
```

Valid Python expression to return true when alert should be "on". Use `value` to access the check result value.

Parameters Add Parameter

Included Entities Event Log Check

```
1 [
2   {
3     "project": "eventlog-data-service-zompy",
4     "environment": "live",
5     "type": "zompy"
6   },
7   {
8     "project": "eventlog-data-service-zompy",
9     "type": "zompy"
10 }
11 ]
```

Sharing and reuse of alerts and checks

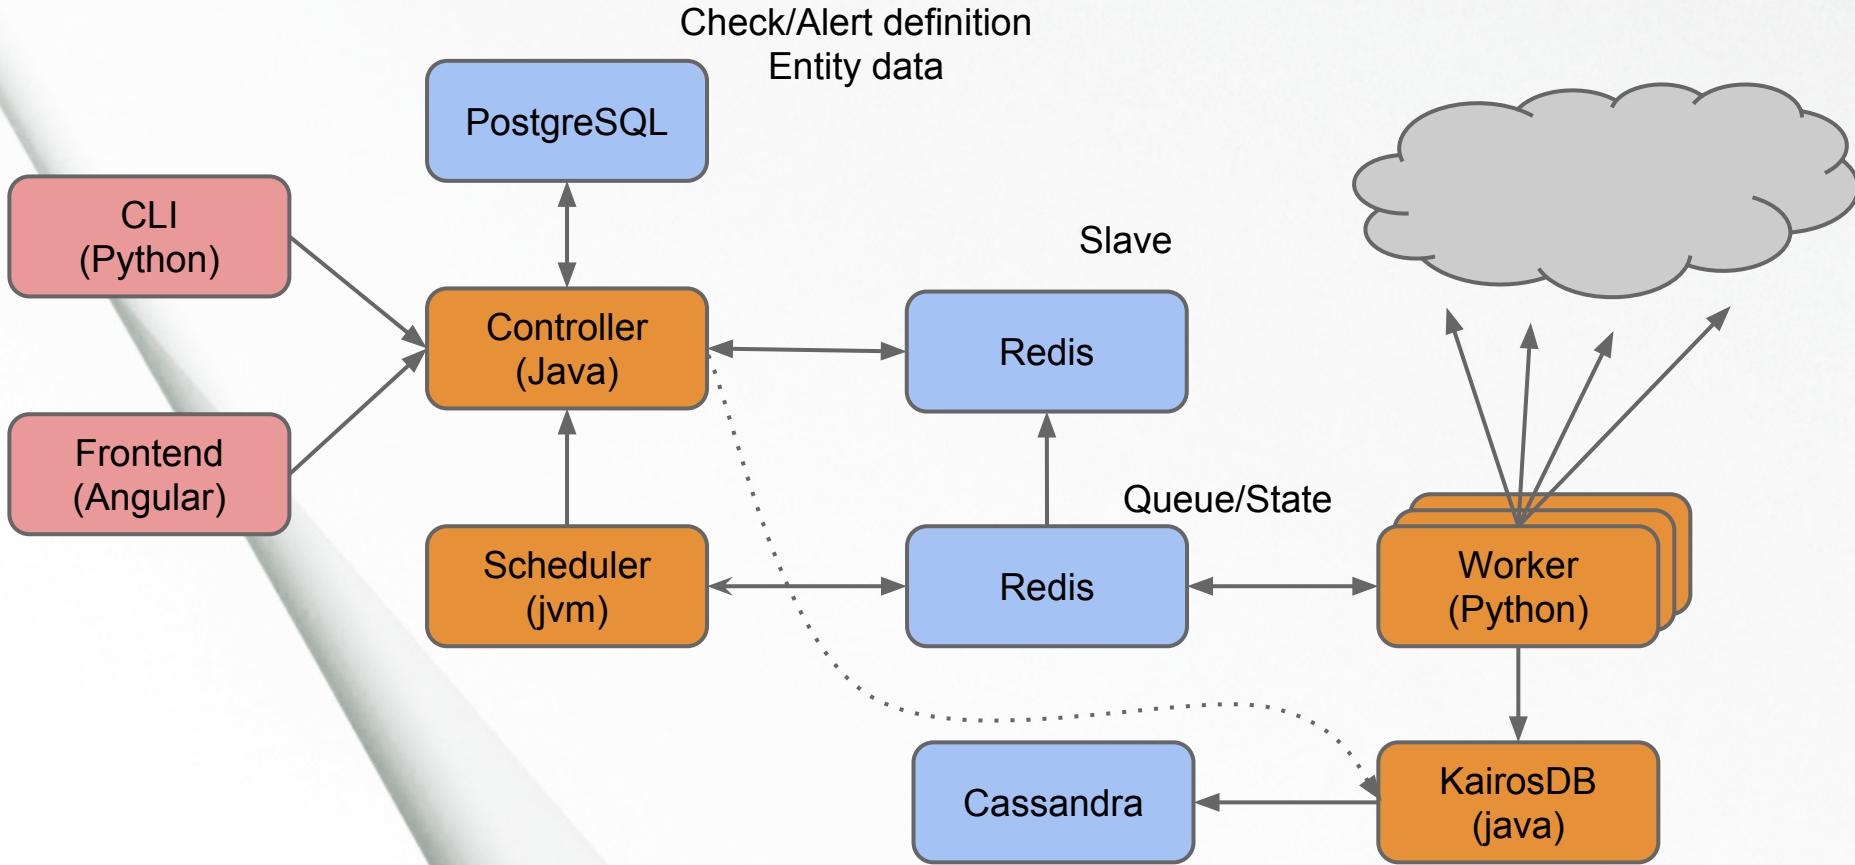
Anyone can add alerts to checks

Alerts are owned by team

Monitor application boundaries/dependencies

Make use of inheritance to customize

ZMON Core + UI + KairosDB



Vagrant Box deploys Docker images

```
→ github-zmon git:(master) vagrant ssh
vagrant@zmon:~$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
fa7a9cf51be9        os-registry.stups.zalan.do/stups/zmon-scheduler-ng:0.1.10   "java -jar zmon-sche"   18 minutes ago    Up 18 minutes
bcae089d51bc        os-registry.stups.zalan.do/stups/zmon-worker:0.1.10      "python app.py"       19 minutes ago    Up 19 minutes
cebb102bfd9a        os-registry.stups.zalan.do/stups/zmon-controller:0.1.9     "catalina.sh run"     19 minutes ago    Up 19 minutes
90053f1ccae7        os-registry.stups.zalan.do/stups/zmon-eventlog-service:0.1.9   "java -jar zmon-even"  21 minutes ago    Up 21 minutes
2b9b36e79cc5        os-registry.stups.zalan.do/stups/zmon-kairosdb:0.1.6      "/usr/bin/config-kai" 22 minutes ago    Up 22 minutes
758ea982e432        abh1nav/cassandra:latest                               "/sbin/my_init"       27 minutes ago    Up 27 minutes
58097a3e0869        os-registry.stups.zalan.do/stups/zmon-redis:0.1.4      "/entrypoint.sh redi" 34 minutes ago    Up 34 minutes
bef7fd8cfe48        os-registry.stups.zalan.do/stups/zmon-ldap:0.1.4      "/slapd.sh"           36 minutes ago    Up 36 minutes
6cd8bc4ae386        os-registry.stups.zalan.do/stups/zmon-postgres:0.1.4     "/docker-entrypoint." 36 minutes ago    Up 36 minutes
```

ZMON 2.0 Dashboards 0 in queue, 2/2 active workers, 0.00/s 15:21 ? Log in

Filter alerts 0 0 0 Example Team 4m

Example Alert

GLOBAL (35.68)

100
50
0
15:18 15:21

Downtimes

- Set or schedule downtimes using the UI
- Use API to automate downtimes, e.g. in deployment tool

Alerts/Checks	Downtimes (2)	History	Children (1)	
Alerts (2)	In downtime (2)	OK (2)		
	Name	Timestamp	Last run	Value
	elsn01:5827	31.08.15 16:01 (6s ago)	33ms	1
	elsn02:5827	31.08.15 16:01 (6s ago)	8ms	2
	elsn03:5827	31.08.15 16:01 (6s ago)	34ms	2

Extendability - Check and Alert functions

- Improve user experience through provided functions

Details	
Check Name, ID, History	Exceptions ZMON Worker in AWS, ID: 1997, History
Check Team	Platform/Software
Check Command	<code>scalyr().errors()</code>
Check Interval	30s
Check Entities	type = application

Extendability - Check and Alert functions

```
def errors(self, minutes=30):
    return self.timeseries(self.default_ts_id, minutes)

def timeseries(self, ts_id, minutes=30, buckets=1):

    val = {
        'token': self.read_key,
        'queries': [
            {'timeseriesId':ts_id,
             'startTime': str(minutes) +'m',
             'buckets': buckets}
        ]
    }

    r = requests.post(self.timeseries_url, data=json.dumps(val), headers={"Content-Type": "application/json"})
    j = r.json()
    if j['status'] == 'success':
        if len(j['results'][0]['values'])==1:
            return j['results'][0]['values'][0] * minutes
        return map(lambda x: x * minutes / buckets, j['results'][0]['values'])
    return j
```

A black and white photograph of the Cliffs of Moher in Ireland. The image shows the massive, layered limestone cliffs that drop sharply into the Atlantic Ocean below. The sky is filled with scattered clouds. In the foreground, there's some grassy vegetation. The text "The Microservices World" is overlaid on the lower half of the image.

The Microservices World

Key Metrics for your service?

- Request rates
- Response rates by HTTP status code
- Latency

Expose your data

```
{  
  "zmon.response.200.GET.checks.all-active-check-definitions.count": 10,  
  "zmon.response.200.GET.checks.all-active-check-definitions.fifteenMinuteRate": 0.18076110580284566,  
  "zmon.response.200.GET.checks.all-active-check-definitions.fiveMinuteRate": 0.1518180485219247,  
  "zmon.response.200.GET.checks.all-active-check-definitions.meanRate": 0.06792011610723951,  
  "zmon.response.200.GET.checks.all-active-check-definitions.oneMinuteRate": 0.10512398137982051,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.75thPercentile": 1173,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.95thPercentile": 1233,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.98thPercentile": 1282,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.999thPercentile": 1282,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.99thPercentile": 1282,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.max": 1282,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.mean": 1170,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.median": 1161,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.min": 1114,  
  "zmon.response.200.GET.checks.all-active-check-definitions.snapshot.stdDev": 42,  
}
```

Start tracking your metrics

Alert: Actuater AWS Metrics

ID: 4953, Status:  , Team: Platform/Software
 Edit  Clone  Inherit  Delete  History  Trial Run  Evaluate  Report an Incident  Comments (0)

Description

.

Condition

False

Responsible Team Platform/Software Details**Check Name, ID, History** REST Metrics (AWS), ID: 2132, History**Check Team** Platform/Software**Check Command**

```
def check():
    j = http(url=':{}//metrics'.format(entity['ports'].values()[-1][0]) if 'ports' in entity else '/metrics',
              oauth2=True).actuator_metrics()

    # workaround for charting if metrics endpoint not covered
    j["metrics"]=[{"GET": {"200": {"mRate": 0, "count": 0, "75th": 0, "99th": 0, "median": 0, "min": 0, "max": 0}}}

    return j
```

Check Interval 1m**Check Entities** type = instance

Display application statistics

kio

Load

Response time

Errors

98.11 req/min	2 instances	2202 ms	2.7%	1.52 4XX/min	1.17 5XX/min	1771 ms	2 instances
------------------	----------------	------------	------	-----------------	-----------------	------------	----------------

zmon data service

Load

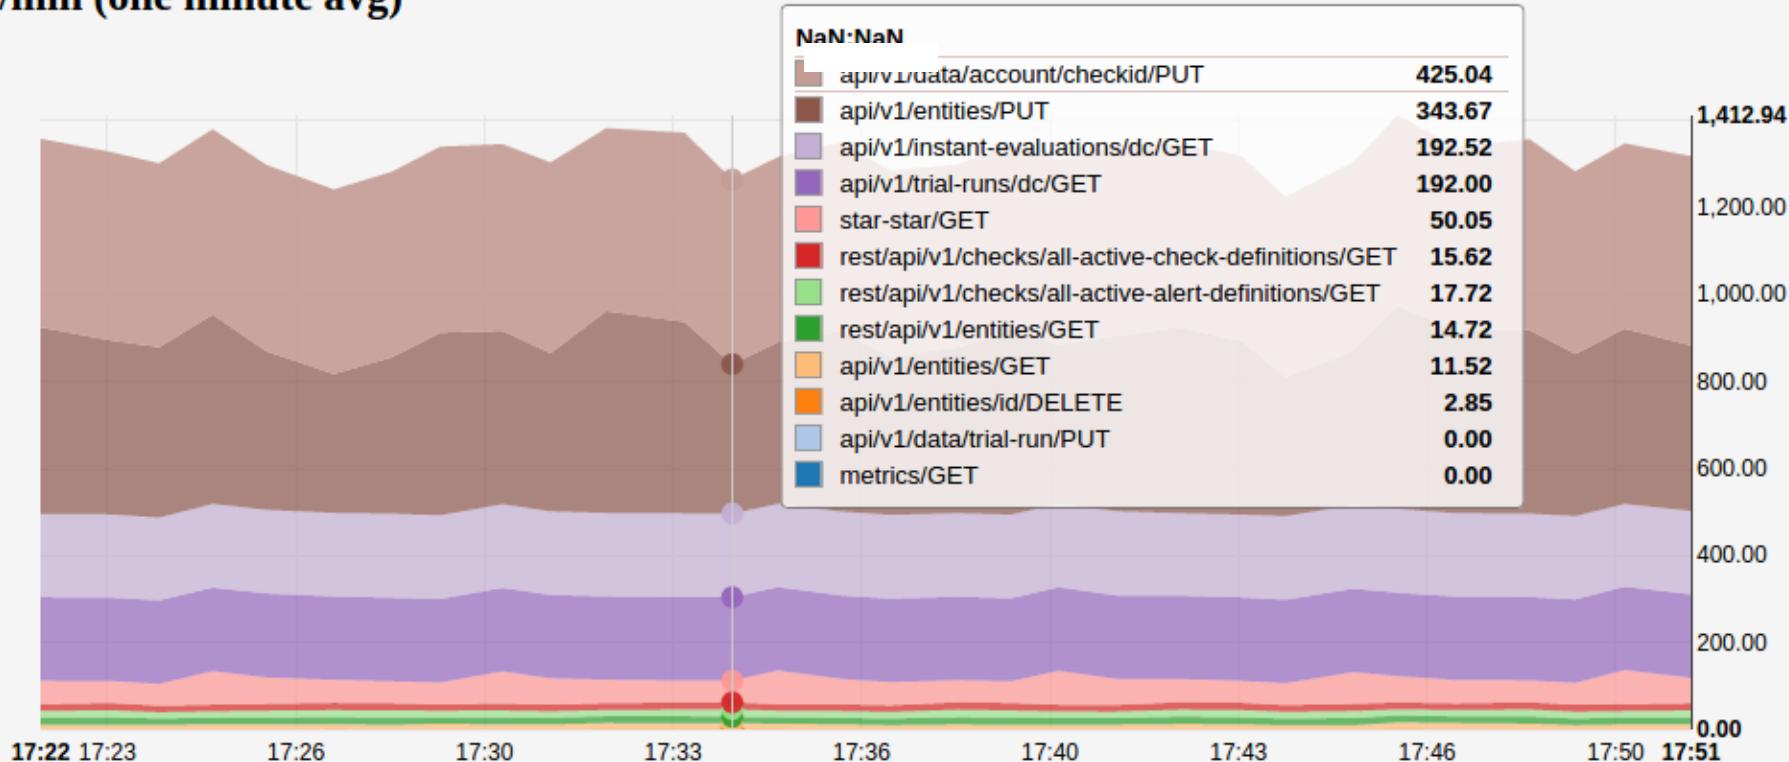
Response time

Errors

1425.34 req/min	2 instances	3508 ms	5.1%	72.98 4XX/min	0.00 5XX/min	19 ms	2 instances
--------------------	----------------	------------	------	------------------	-----------------	----------	----------------

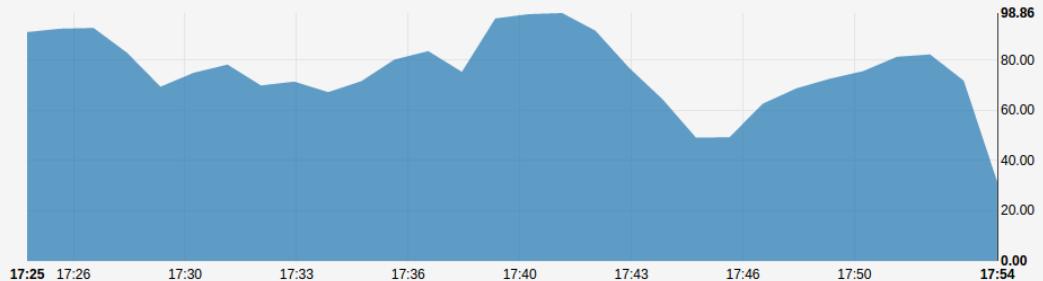
Application metrics

rate/min (one minute avg)

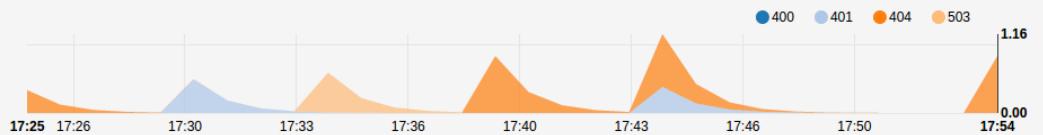


Request rates for status 200

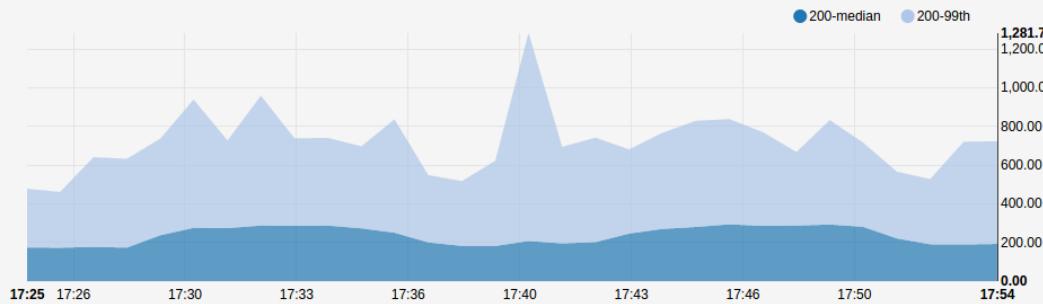
Continued ...



Request rates other



Response times for status 200



Reuse of check

Alert: Unauthorized Access to Event Data Service {rate}

ID: 5255, Status:  , Team: STUPS

 Edit  Clone  Inherit  Delete  History  Trial Run  Evaluate  Report an Incident  Comments (0)

Description

Monitoring 401 on EventLog Data Service Endpoint

Condition

```
def alert():
    rate = 0
    for k,v in value.iteritems():
        for hm, vs in v.iteritems():
            ks = vs.keys()
            if '401' in ks:
                rate += vs['401']['mRate']
    return capture(rate=rate) > 0.1
```

Responsible Team STUPS

Details

Check Name, ID, History REST Metrics, ID: 2115, [History](#)

Check Team Platform/Software

Libraries available for

Spring boot

<https://github.com/zalando/zmon-actuator>

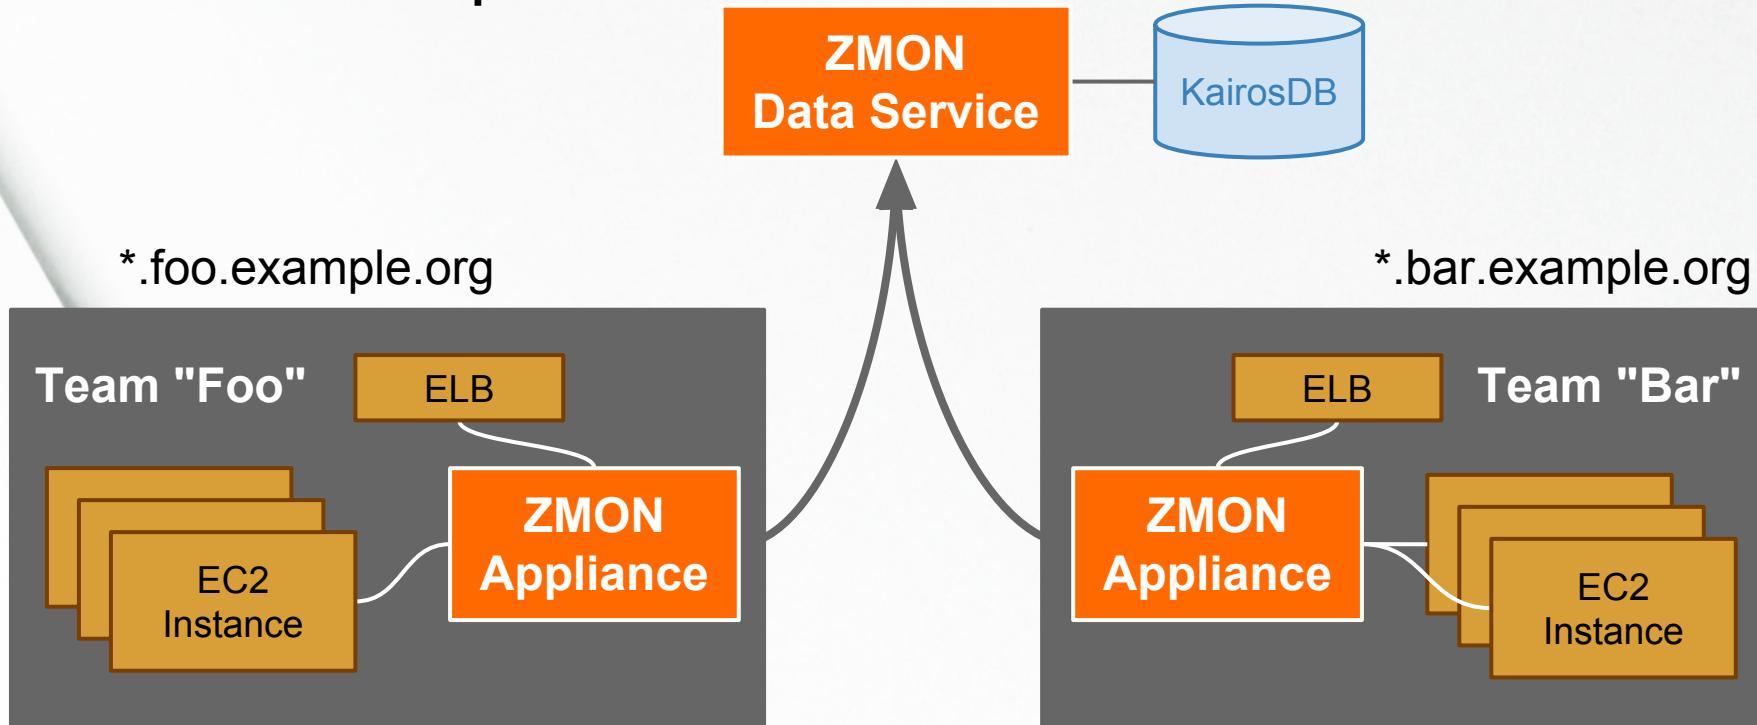
Clojure

<https://github.com/zalando-stups/friboo/>

Play (done, to be released)

Multi DC / AWS

ZMON in AWS Setup



Multi DC / Zone deployment possible

- Scheduler supports queue filters by entity
 - e.g. {"dc":"dc1"} vs {"dc":"dc2"} queue filters
- Scheduler can apply base filter
 - only handles entities with {"dc":"dc1"}
- Worker can report home using:
 - Redis (we use this across DCs)
 - HTTPS (AWS->DC)

ZMON AWS Agent

Uses Amazon API to fetch:

- ELBs
- EC2 instances
- RDS instances

Pushes enriched entities to entity service

Prometheus?

read "text" result

Kubernetes Example: Exports in Prometheus text format

```
kubelet_docker_operations_latency_microseconds{operation_type="inspect_container",quantile="0.9"} 9602
kubelet_docker_operations_latency_microseconds{operation_type="list_containers",quantile="0.9"} 9740
```

Check Command

```
1 def check():
2     d = http(':10255/metrics').text()
3     r= []
4     for l in d.split("\n"):
5         if l.startswith("kubelet_docker_operations_latency_microseconds"):
6             lp = l.split(" ")
7             rv=' '.join(lp[1:])
8             if "NaN"==rv:
9                 continue
10            a=lp[0]
11            a = a[a.find('{')+1:a.rfind('}') ]
12            a = [(v.split('=')[0], v.split('=')[1]) for v in a.split(",")]
13            c = r
14            while True:
15                b=a.pop(0)[1].replace("'", '')
16                if len(a)==0:
17                    c[b]=rv.replace("'", '')
18                    break
19                else:
20                    if b in c:
```

Valid Python expression used to query the instances specified in entities filter.

Yields a usable nested dictionary

```
{"list_images":  
    {"0.9":"120252",  
     "0.99":"120252",  
     "0.5":"120252"},  
"version":  
    {"0.9":"1281",  
     "0.99":"2183",  
     "0.5":"873"},  
"list_containers":  
    {"0.9":"9740",  
     "0.99":"23378",  
     "0.5":"3717"},  
"inspect_container":  
    {"0.9":"9602",  
     "0.99":"18367",  
     "0.5":"4419"}  
}
```

Internals

ZMON's basic data flow



```
{ "check":  
  { "id": 1,  
    "entity": { "host": "monitor01" },  
    "command": "snmp().load()",  
    "alerts": [  
      { "id": 100,  
        "condition": "value['load1'] > 10" }  
    ]  
  }  
}
```

ZMON's basic data flow



```
-- store check result "snmp().load()"  
lpush zmon:check:1:monitor01 {"load1":5,"load5":3,"load15":2}  
  
-- keep last 20 results (for dashboard charts)  
ltrim zmon:check:1:monitor01 20  
  
-- alert active?  
sadd zmon:alert:100 monitor01  
  
-- alert inactive?  
srem zmon:alert:100 monitor01
```

ZMON Vagrant Box:

<https://github.com/zalando/zmon>

ZMON Homepage:

<https://zalando.github.io/zmon>

Zalando Tech:

<https://tech.zalando.com>

