

Revisiting the Outbox Pattern in Go

Nikolay Kuznetsov

10 Sep 2025

SEP 09 – 11, 2025
CONTAINER
days
CONFERENCE



About me

Senior Software Engineer

Zalando Helsinki 

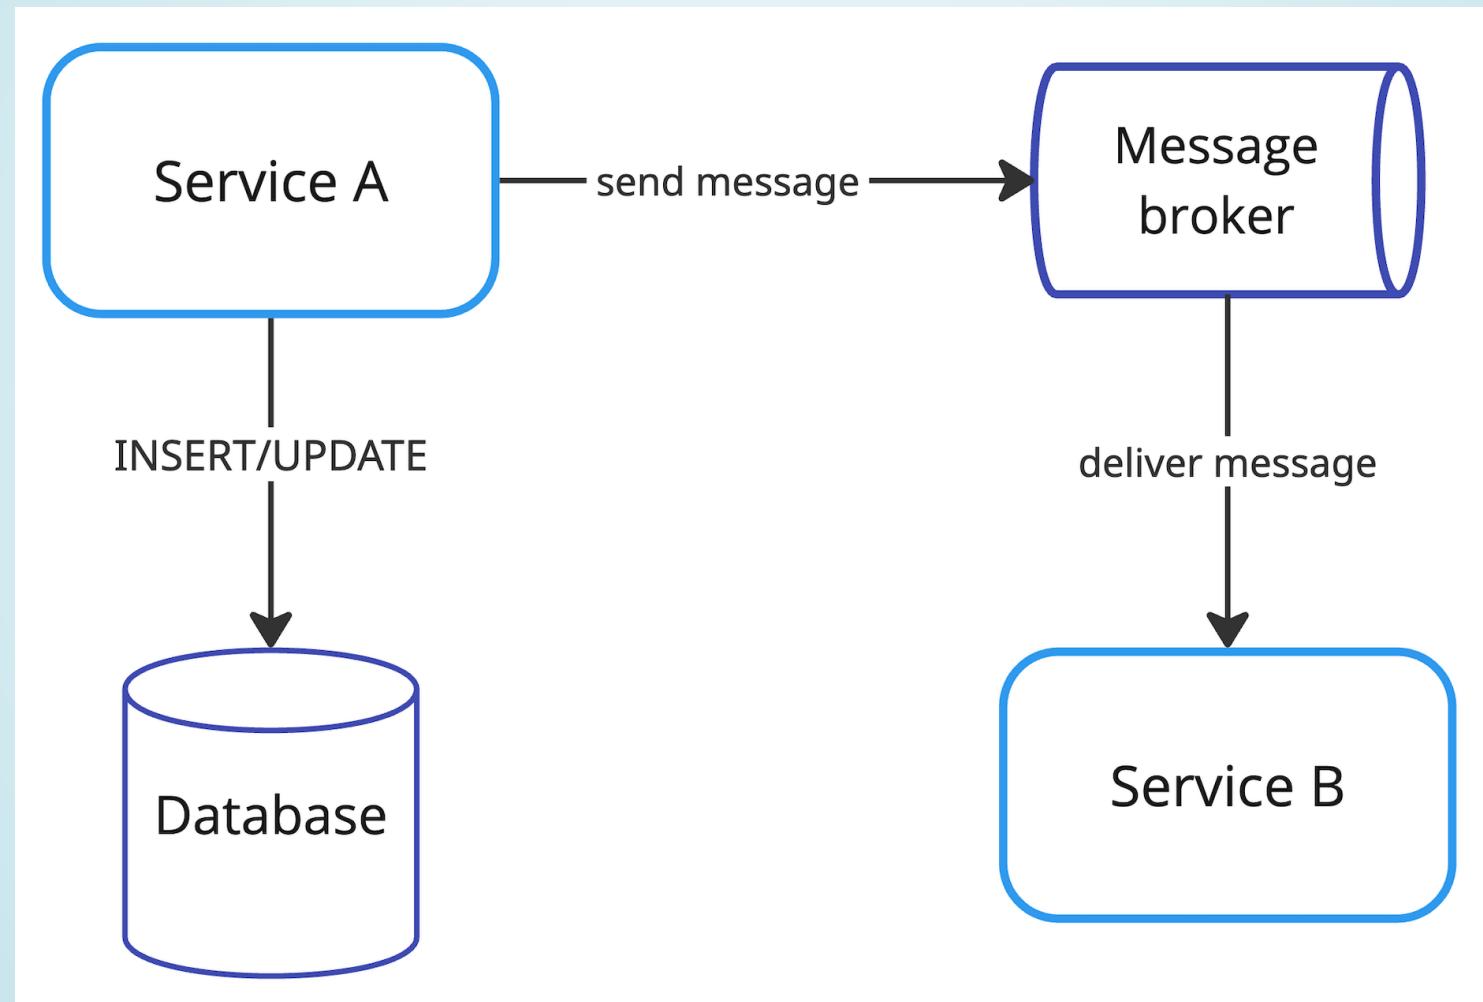
C → Java → Kotlin → Go

Author of [pgx-outbox](#) library

Talk inspirations

- Revisiting the Outbox Pattern
 - article by Gunnar Morling
 - Personal experience

Dual write problem



Dual write problem

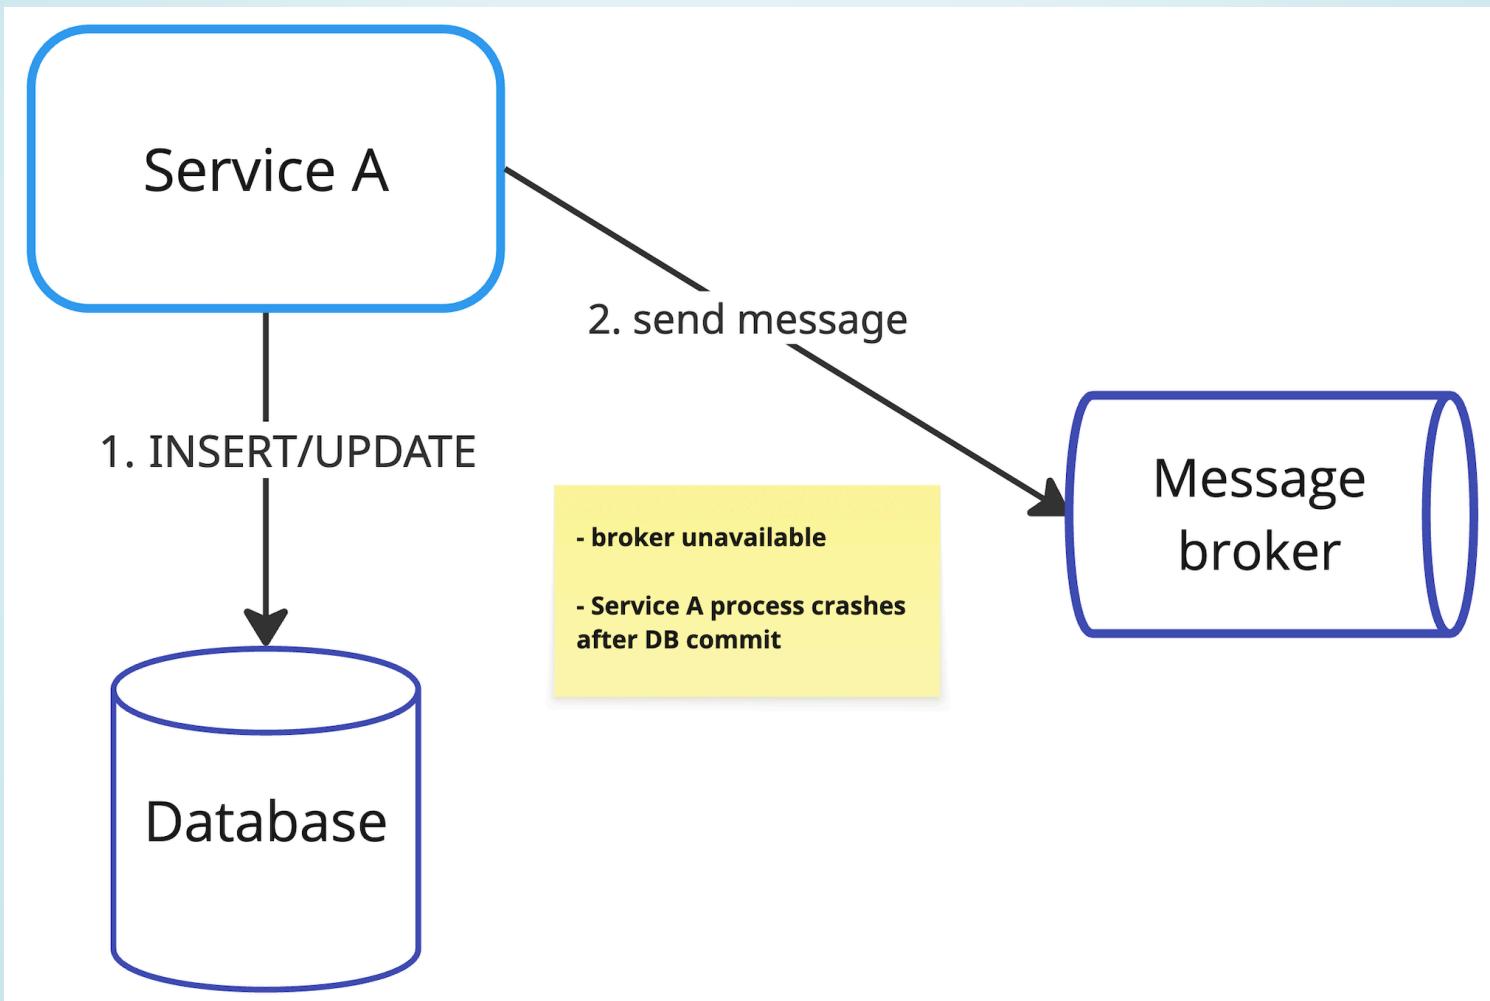
Service A writes to DB and message broker

Inconsistent system state if any of writes fails

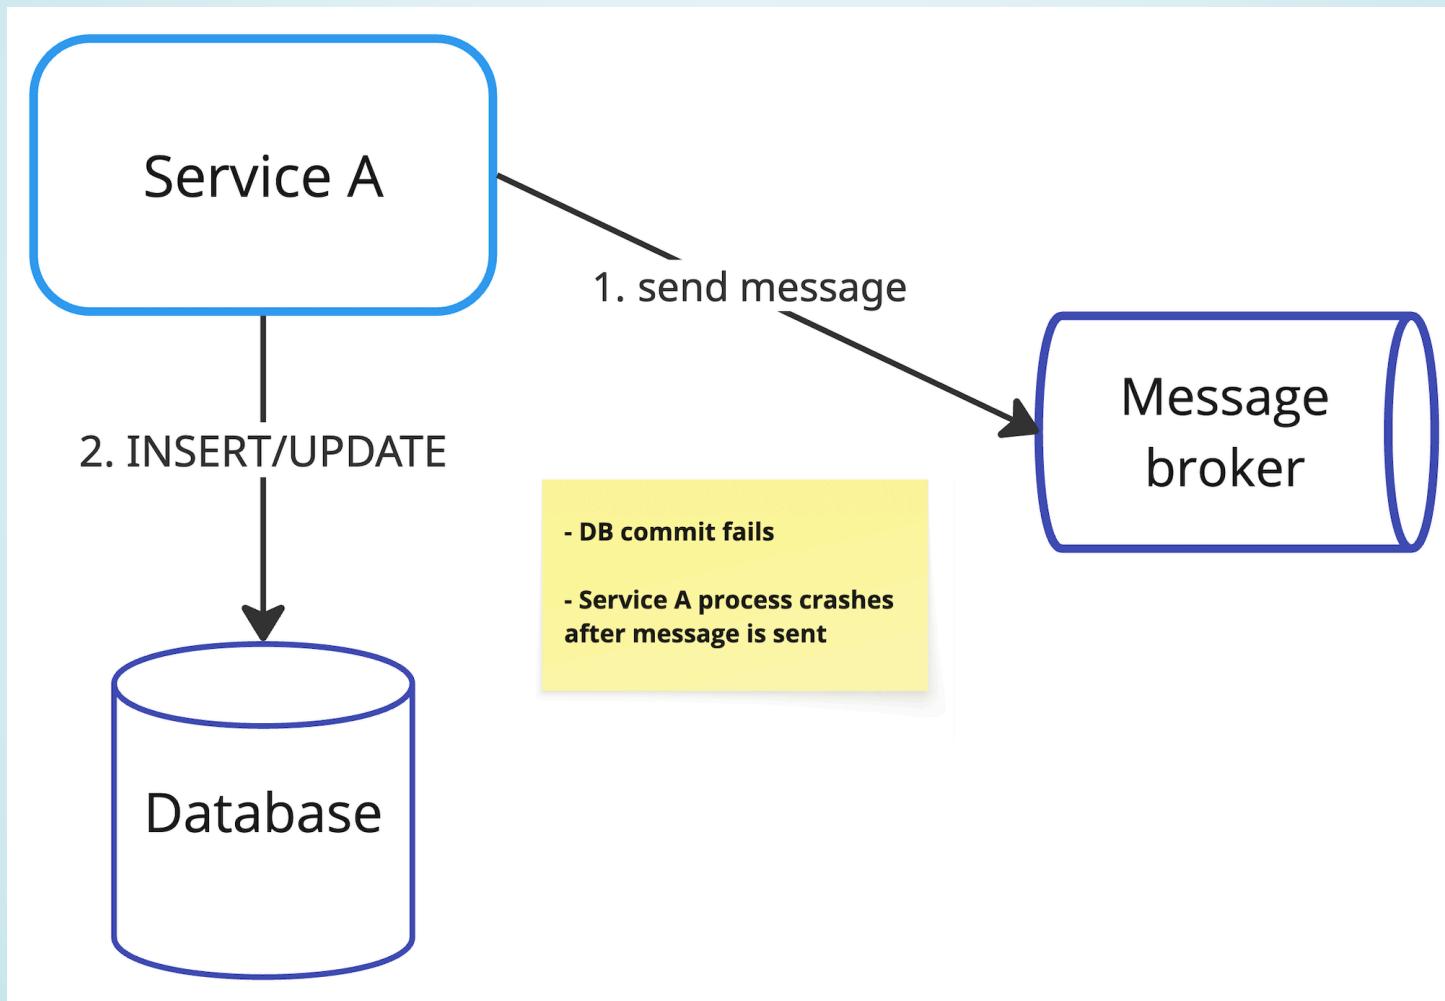
Naive approaches

- Write to DB first, then send to broker
- Send to broker first, then write to DB
- Send to broker within DB tx before commit

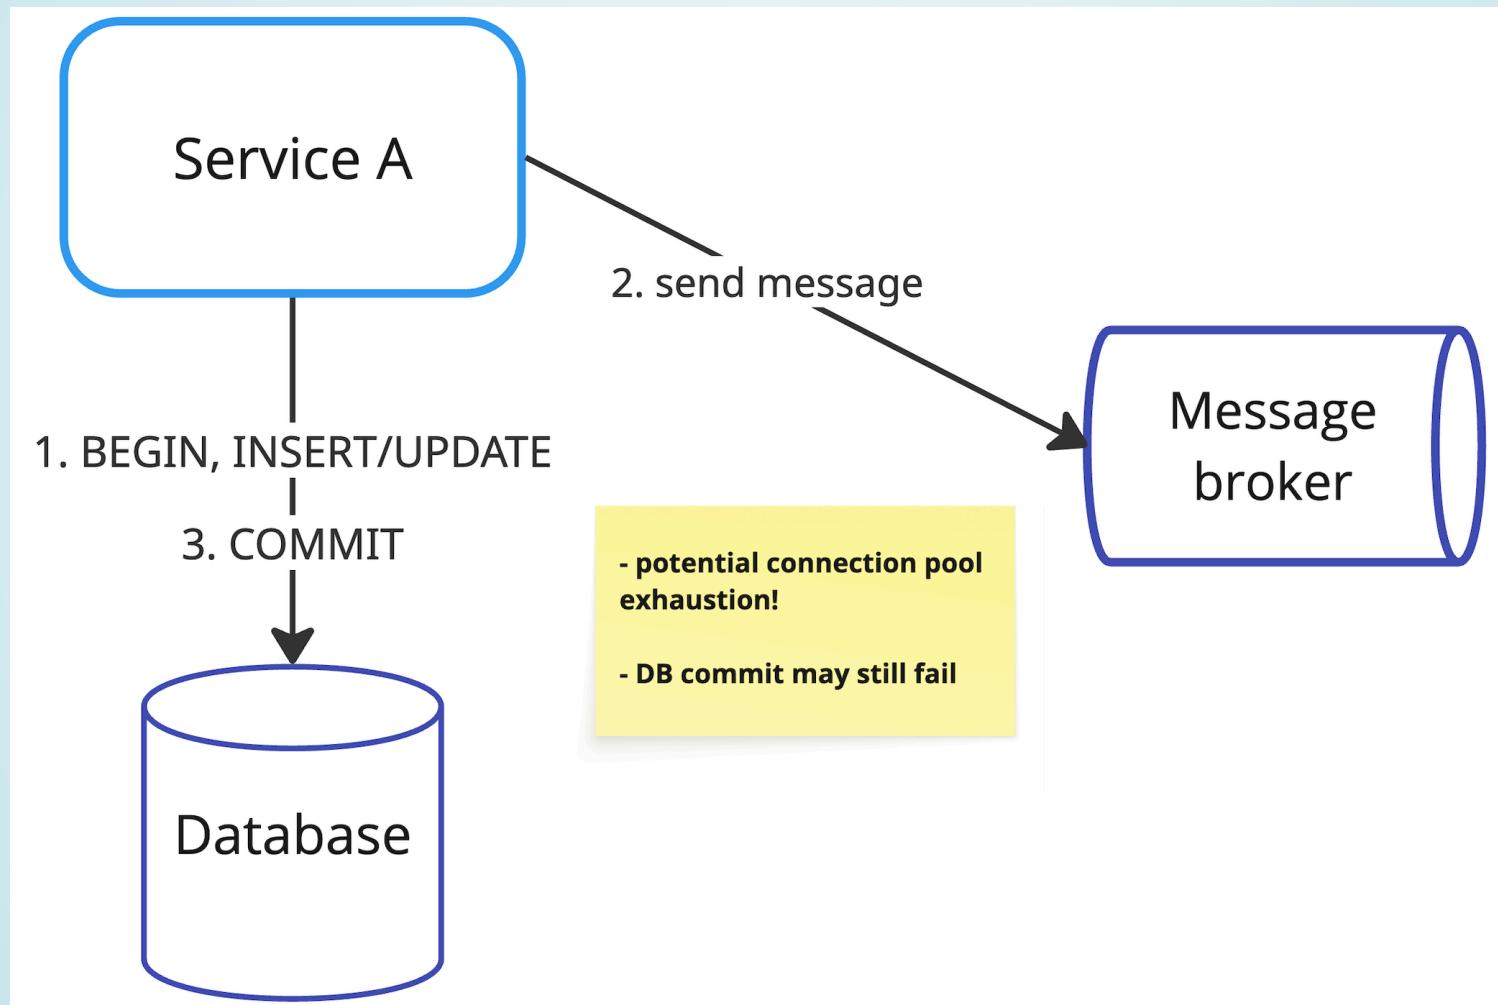
Database first



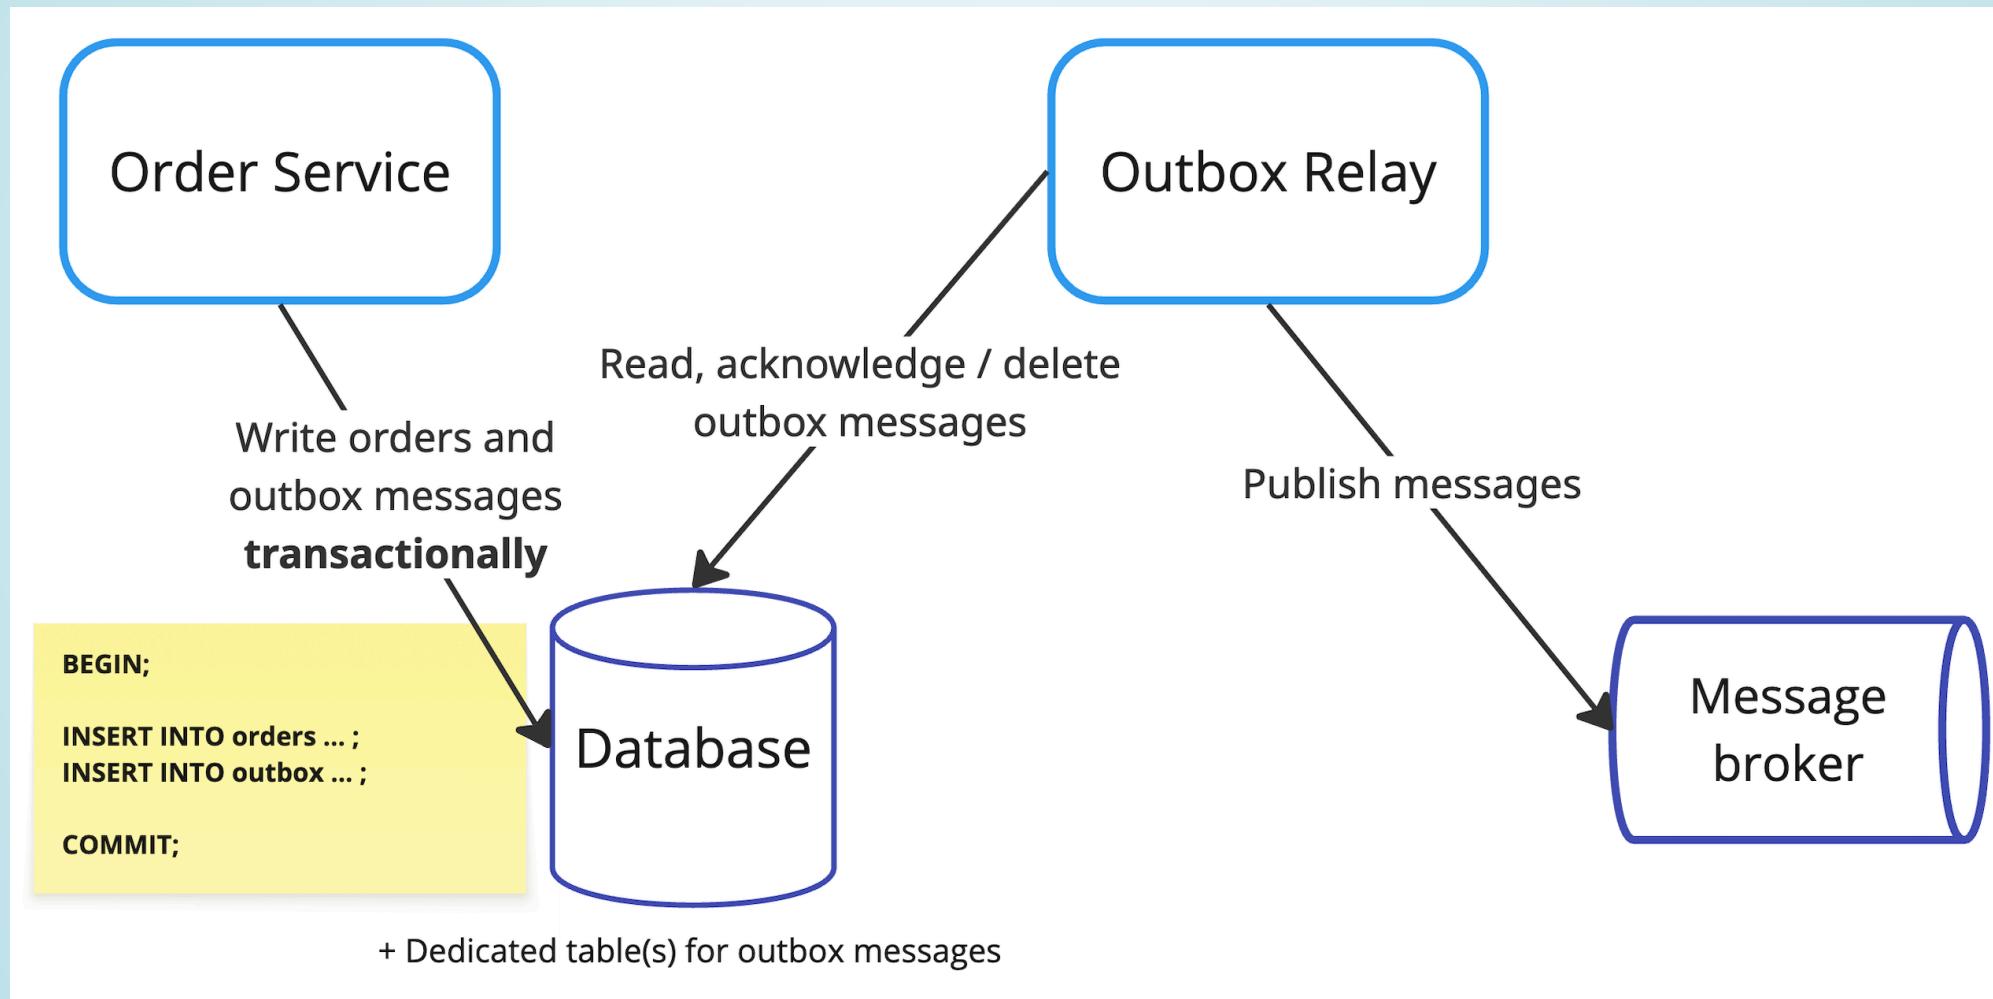
Broker first



Send inside DB tx



Transactional outbox pattern



Outbox table example

```
CREATE TABLE outbox (

    id          BIGINT PRIMARY KEY GENERATED ALWAYS AS IDENTITY,
    payload     JSONB                      NOT NULL,
    created_at  TIMESTAMP DEFAULT CURRENT_TIMESTAMP NOT NULL,
    -- additional metadata: event type, domain object type and id, topic
);
```

Outbox pattern pros

one of solutions to dual write problem

- + *Atomicity* between orders and outbox messages
- + At least once delivery to message broker

Outbox pattern cons

- Message publication delay
- Extra code or deployment unit
- Consumer might need to deduplicate
 - Increased database load

Reading from outbox

id	payload	created...	broker	topic
91	{"id": "f20ea629"}	2025-08-17 14:10...	sns	topic1
90	{"id": "3669a2c5"}	2025-08-17 14:10...	sns	topic1

- Periodic polling
- Log-based change data capture (CDC)

Periodic polling

- SELECT (with LIMIT)
- Options to mark published:
 - a. DELETE by ids
 - b. UPDATE *published_at* by ids
 - c. Store latest id in a separate table

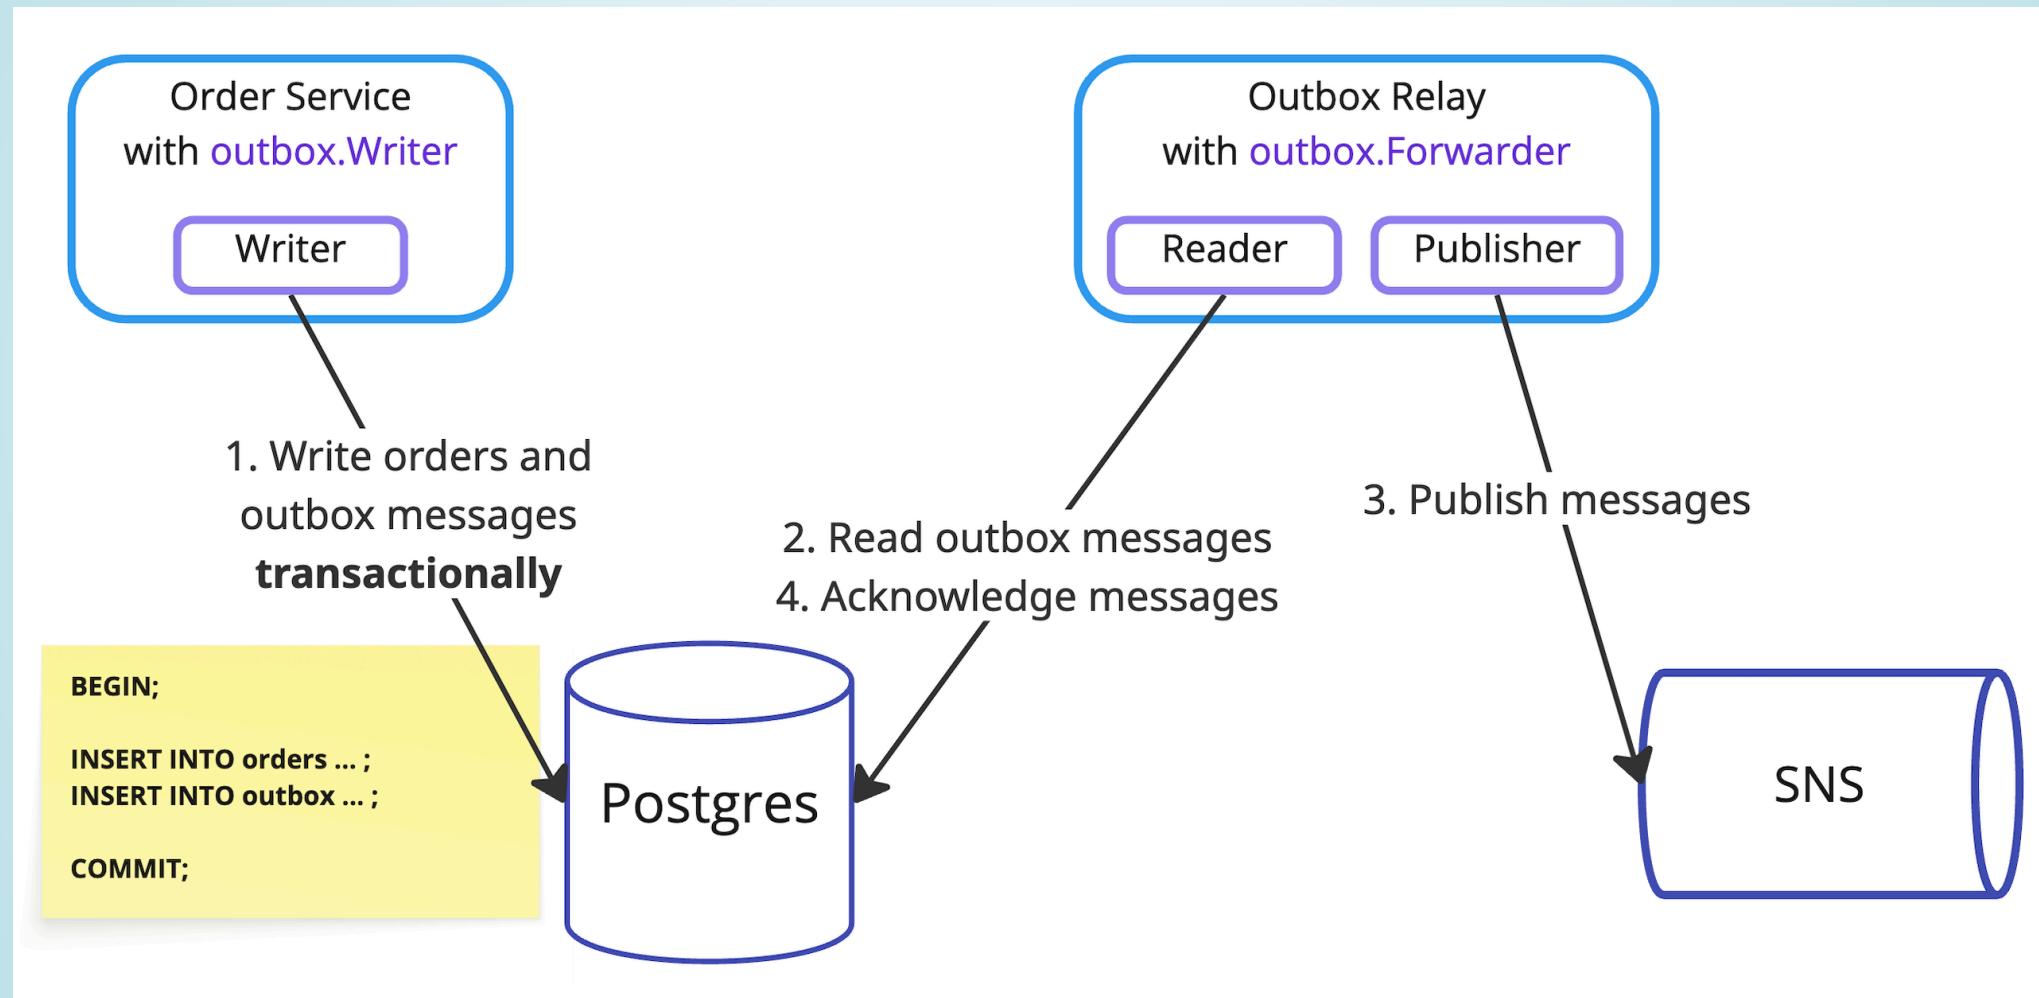
Go libs with polling

- Watermill-SQL
 - pgx-outbox
 - go-outbox, ▪ outboxer
- general purpose queues: ▪ River, etc

pgx-outbox

- Strong focus on *Postgres* and *pgx* driver
 - Implements *Message*, *Writer*, *Forwarder*
- UPDATE approach to mark published messages
 - Inspired by Watermill Forwarder

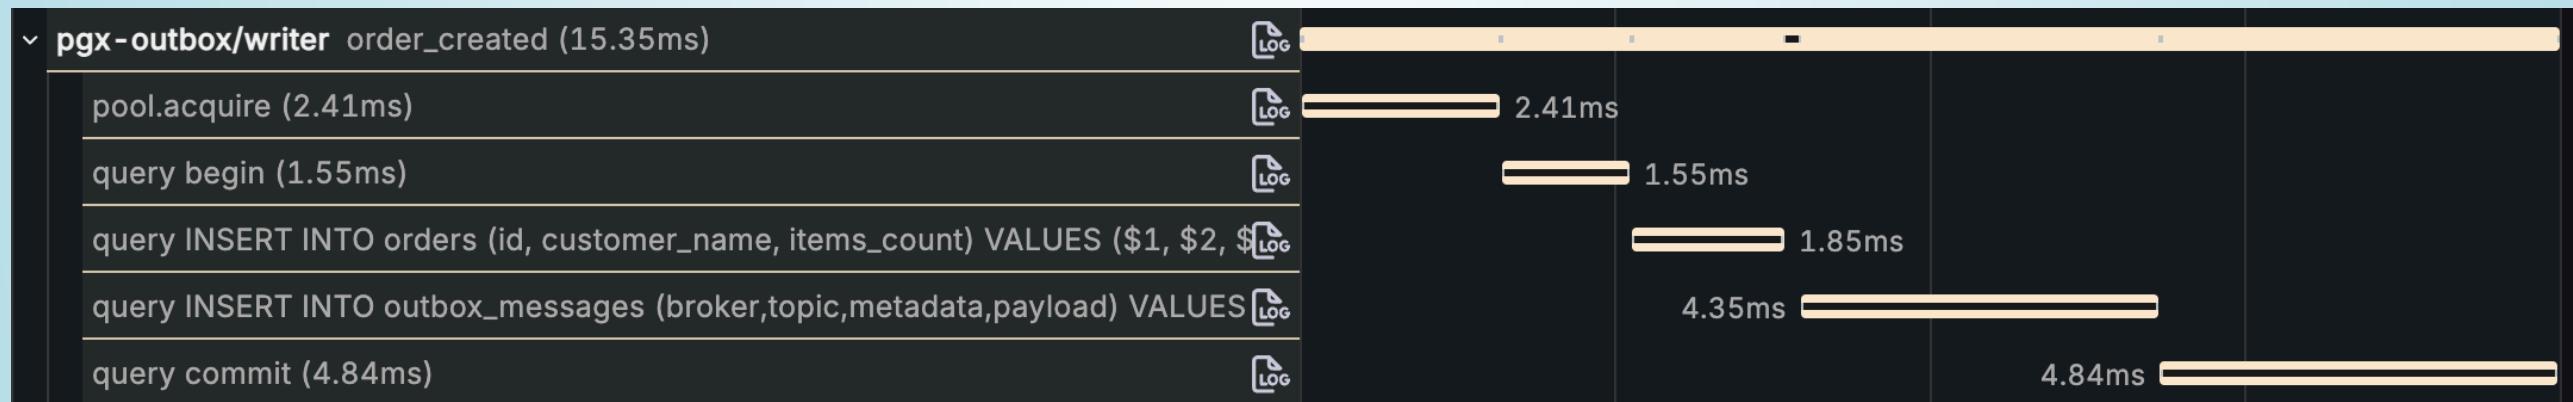
pgx-outbox architecture



Outbox Writer

```
type Writer interface {  
    // supports both pgx.Tx and stdlib *sql.Tx  
    Write(ctx, tx Tx, message Message) (int64, error)  
  
    // pgx transaction only to invoke SendBatch and Prepare methods  
    WriteBatch(ctx, tx pgx.Tx, messages []Message) ([]int64, error)  
}
```

Tracing with otelpgx



Polling cons

- Low publication delay vs polling too often
- Risk of starving other workloads under high traffic
- UPDATE and DELETE create dead tuples (MVCC)
- Autovacuum settings might need tuning

Log-based CDC

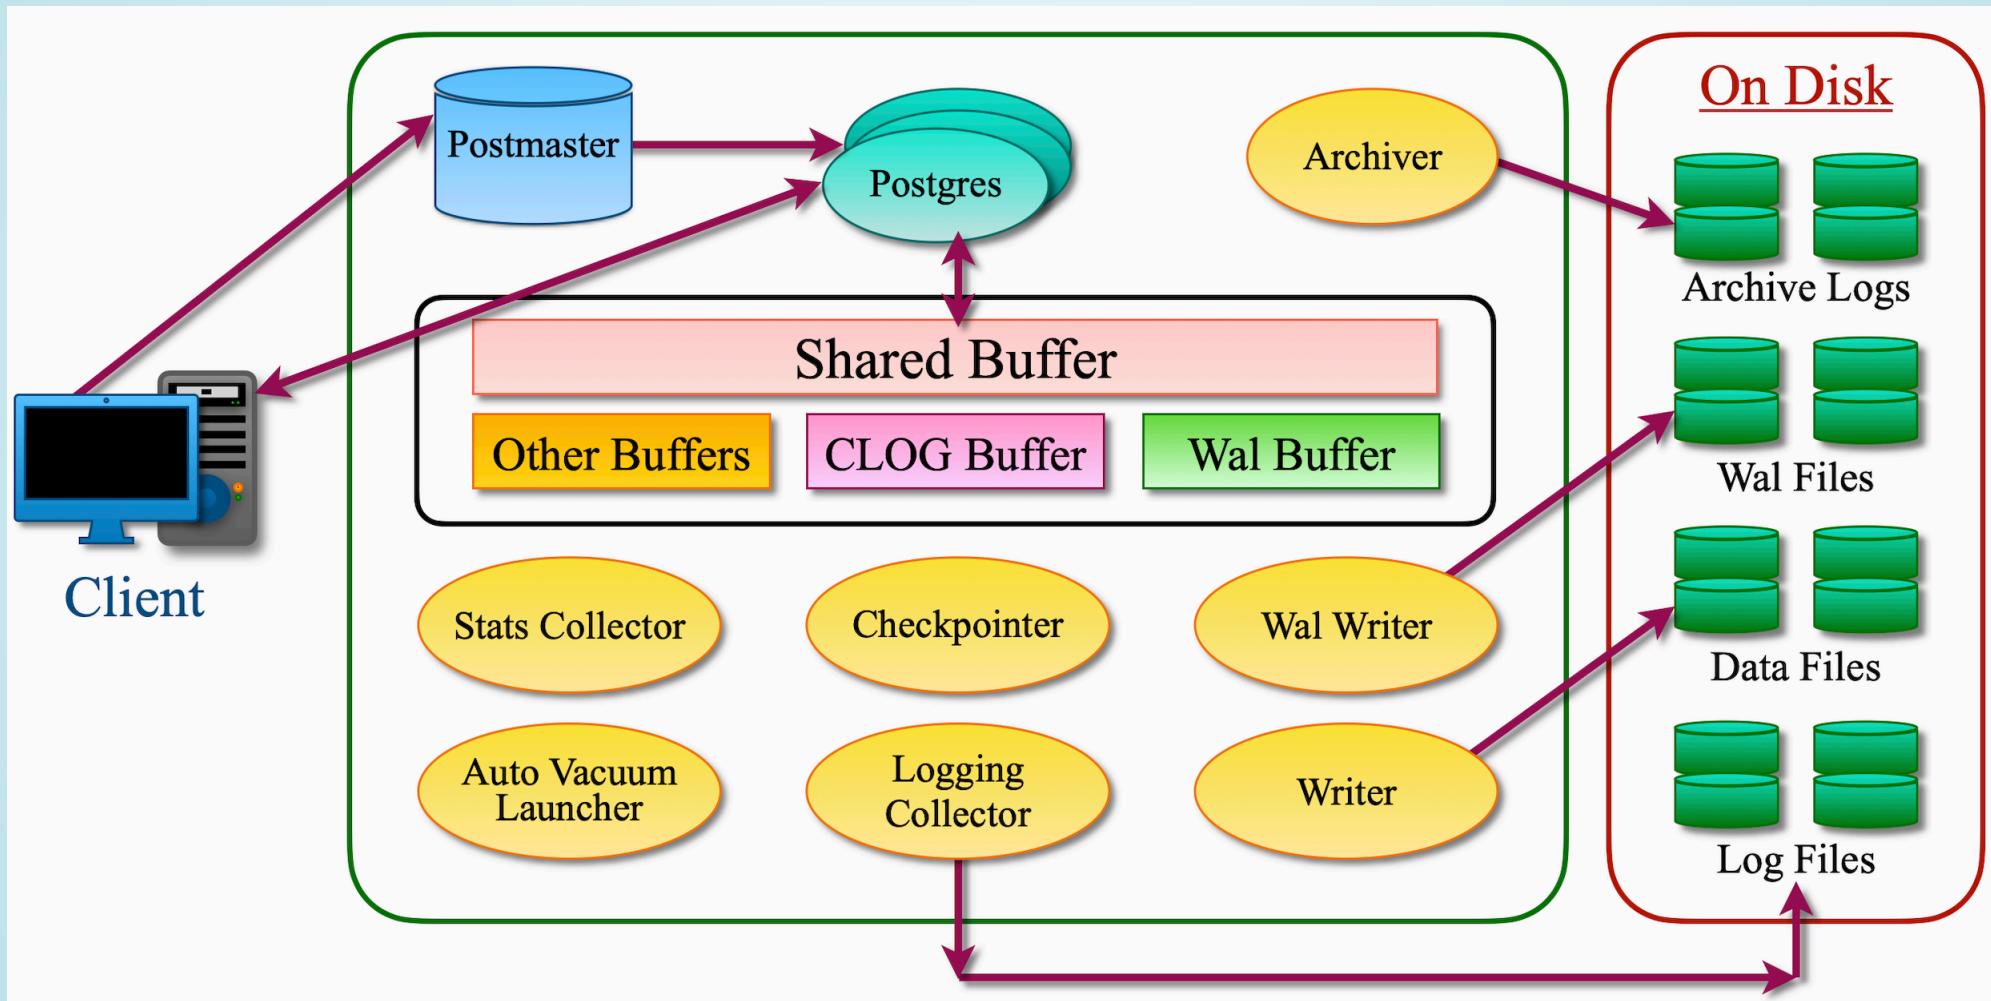
Write-Ahead Log (WAL) in Postgres

- + Lower publication delay (tens of ms)
- + Reduced DB load: avoids polling, dead tuples
- + Events emitted in exact commit order

Write-Ahead Log

- append-only **binary** log of all **row-level** changes
- written before data files, ensures crash recovery
- changes grouped per transaction in commit order
 - exists in shared buffer, flushed to on-disk files

Postgres process and memory architecture



Transactions example

TX1	START TRANSACTION	0/AC90FD
TX1	INSERT INTO t1 VALUES ...	0/AC9102
TX2	START TRANSACTION	0/AC9110
TX2	INSERT INTO t1 VALUES ...	0/AC912F
TX2	UPDATE t1 SET ...	0/AC9134
TX3	START TRANSACTION	0/AC9140
TX3	INSERT INTO t1 VALUES ...	0/AC914A
TX1	UPDATE t1 SET ...	0/AC9153
TX2	DELETE FROM t1 WHERE ...	0/AC9155
TX3	UPDATE t1 SET ...	0/AC915D
TX2	ROLLBACK	0/AC9160
TX3	DELETE FROM t1 WHERE ...	0/AC9168
TX3	COMMIT	0/AC916A
TX1	DELETE FROM t1 WHERE ...	0/AC9170
TX1	COMMIT	0/AC9179

<https://www.dolthub.com/blog/2024-03-08-postgres-logical-replication/>

Corresponding WAL

0/AC9140	BEGIN	TX3
0/AC914A	INSERT t1 tuple [...]	TX3
0/AC914F	INSERT t1 tuple [...]	TX3
0/AC915D	UPDATE t1 old [...] new [...]	TX3
0/AC9160	UPDATE t1 old [...] new [...]	TX3
0/AC9164	UPDATE t1 old [...] new [...]	TX3
0/AC9168	DELETE t1 tuple [...]	TX3
0/AC916A	COMMIT	TX3
0/AC9179	BEGIN	TX1
0/AC9102	INSERT t1 tuple [...]	TX1
0/AC9153	UPDATE t1 old [...] new [...]	TX1
0/AC9170	DELETE t1 tuple [...]	TX1
0/AC9176	DELETE t1 tuple [...]	TX1
0/AC9179	COMMIT	TX1

<https://www.dolthub.com/blog/2024-03-08-postgres-logical-replication/>

Reading from WAL

Logical replication protocol in Postgres

- Debezium (Server)
 - *jackc/pglogrepl*
- PeerDB

Debezium

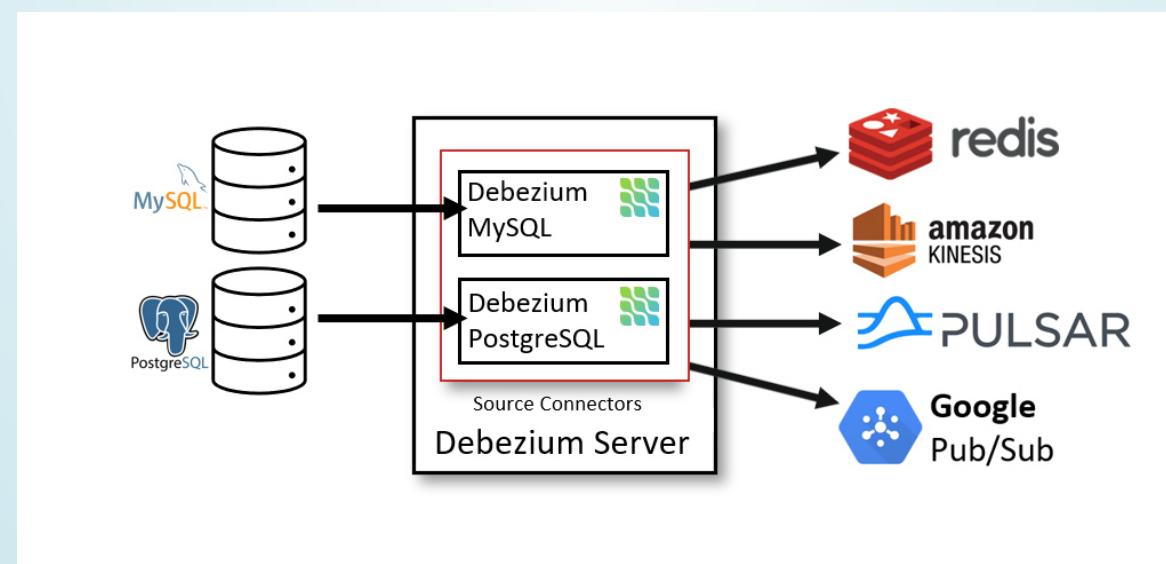
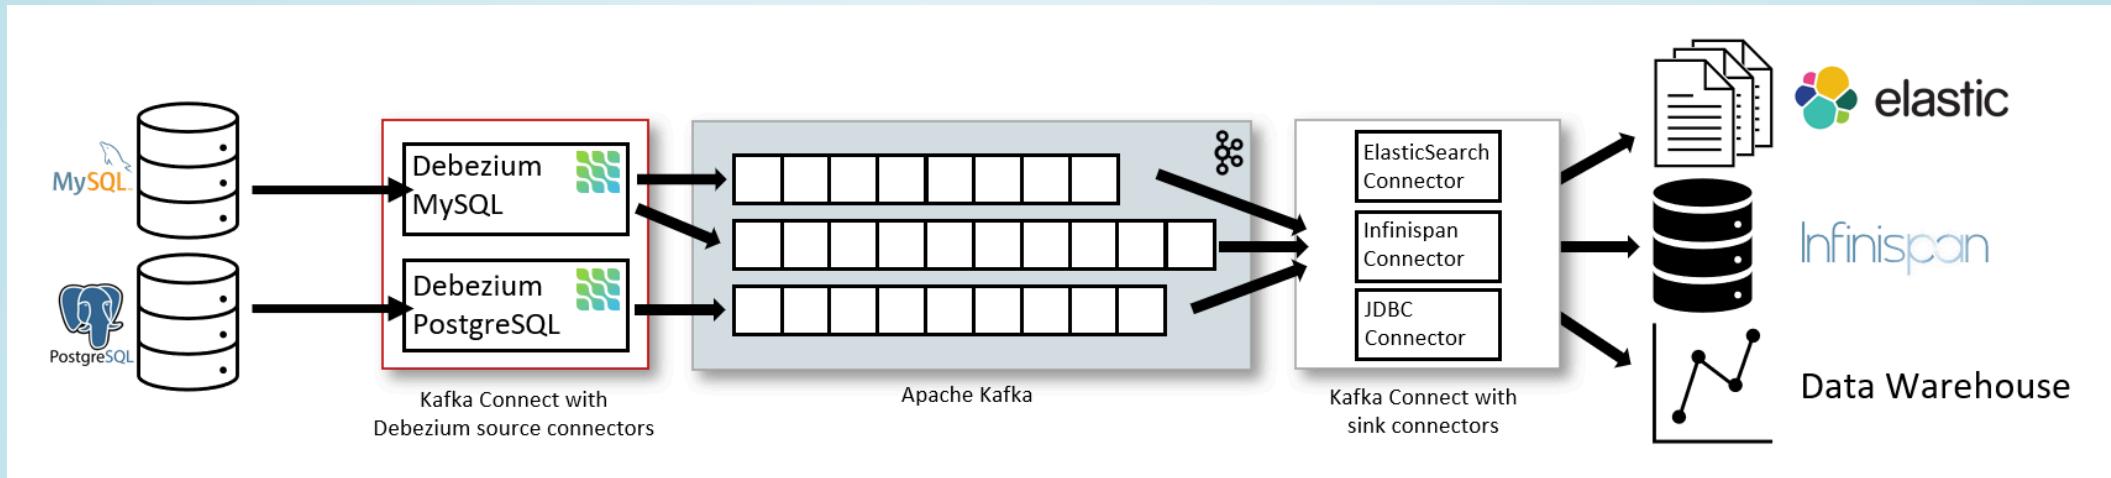
Connectors for Kafka Connect or standalone Server

Consumes DB events and can send to sinks

Very mature, started in 2015 by Red Hat

Open source, written in **Java**

Debezium architectures



<https://debezium.io/documentation/reference/stable/architecture.html>

Server config example

```
# Source: Postgres
debezium.source.connector.class=io.debezium...PostgresConnector
debezium.source.database.{}=
debezium.source.plugin.name=pgoutput
debezium.source.publication.name=debezium_publication
debezium.source.slot.name=debezium_slot
debezium.source.table.include.list=public.outbox_messages

# Sink: HTTP
debezium.sink.type=http
debezium.sink.http.url=http://http-listener:8080

# Output format & single message transformation (SMT)
debezium.format.{key,value}=json
debezium.transforms=unwrap
```

Outbox event router

- Implemented as Kafka Connect SMT
 - Suggests outbox table schema:
id, aggregateid, aggregatetype, type, payload
- Routes to Kafka topic based on *aggregatetype*
 - Uses *aggregateid* as message key

Debezium + Go

- Kubernetes CRD + Operator for CDC
- Operator creates single-pod Deployment
 - Pod is powered by Debezium,
captures outbox inserts in *source* DB,
publishes to message broker sink

jackc/pglogrepl

decodes logical replication messages

includes a basic example

400 stars, 200 users

same author as *pgx* driver

not so active development

Concepts

Publication

Replication slot

Decoding plugin

Log sequence number (LSN)

Publication

Definition of tables and CDC operations to be published to consumers

```
CREATE PUBLICATION pub1 FOR TABLE t1, t2;  
-- ALL TABLES  
-- WITH (publish = 'insert, update', 'delete', 'truncate');  
  
SELECT * FROM pg_publication;
```

```
query := "CREATE PUBLICATION pub1 FOR TABLE t1, t2"  
  
result := conn.Exec(ctx, query) // low level *pgconn.PgConn  
defer result.Close()  
  
, err := result.ReadAll()
```

Replication slot

Named position in WAL to track a replication consumer progress

```
SELECT pg_create_logical_replication_slot('slot', 'pgoutput');  
SELECT * FROM pg_replication_slots;
```

```
pglogrepl.CreateReplicationSlot(ctx, conn, "slot", "pgoutput")
```

Decoding plugins

pgoutput - built-in, default

test_decoding - built-in, only for testing

wal2json - 3rd party, produces JSON
decoderbufs, pglogical

pgoutput arguments

```
pluginArguments := []string{
    "proto_version '2'", // versions 3 and 4 are not supported by the lib
    fmt.Sprintf("publication_names '%s'", "pub1"),
    "messages 'false'", // pg_logical_emit_message() is not used
    "streaming 'false'", // receive only committed transactions
}
```

Starting replication

LSN is a unique 64-bit address in WAL, e.g. 0/AC90FD

```
START_REPLICATION SLOT slot LOGICAL 0/AC90FD -- LSN where to start
  ("proto_version" '2',
   "publication_names" 'pub1',
   "messages" 'false',
   "streaming" 'false');
```

```
sysIdent, err := pglogrepl.IdentifySystem(ctx, conn)
pglogrepl.StartReplication(ctx, conn, "slot", sysIdent.XLogPos,
  pglogrepl.StartReplicationOptions{PluginArgs: pluginArguments})
```

Message loop

```
for {
    var rawMsg pgproto3.BackendMessage
    rawMsg, err := conn.ReceiveMessage(ctx)

    switch msg.Data[0] {
        // XLog is historical name for WAL
        case pglogrepl.XLogDataByteID: // 'w'
            err = handleXLogData(msg.Data[1:]) // actual data

        case pglogrepl.PrimaryKeepaliveMessageByteID: // 'k'
            err = handlePrimaryKeepalive(msg.Data[1:])
    }
}
```

Handling data

- Insert messages
- Relation messages
- Status updates:

write, flush, apply positions

Logical replication is hard !



PeerDB

- efficient data streaming from DBs
 - written in Go (and Rust)
- uses *Temporal* orchestration engine
- open-source, Elastic License 2.0 (ELv2)
- acquired by ClickHouse, integrated to ClickPipes

Peers

connection definition to a data source/sink

The screenshot shows the PeerDB application interface. On the left is a sidebar with the following navigation items:

- Peers (selected)
- Mirrors
- Alerts
- Scripts
- Logs
- Settings

The main content area has a title "Peers" and a descriptive paragraph: "Peers represent a data store. Once you have a couple of peers, you can start moving data between them through mirrors." Below this is a table with two rows, each representing a peer.

Peer Name	Peer Type
postgres_source	PostgreSQL
kafka_sink	Kafka

Mirrors

replication flow that continuously syncs data

The image shows a screenshot of the PeerDB web interface. On the left is a sidebar with the PeerDB logo and navigation links: Peers, Mirrors (which is the active page, indicated by a blue background), Alerts, Scripts, Logs, and Settings. The main content area has a title "Mirrors" and a description: "Mirrors are used to replicate data from one peer to another. PeerDB supports three modes of replication. Begin moving data in minutes by following the simple [PeerDB Quickstart](#)". Below this is a section titled "Change-data capture" with a table showing a replication setup:

Name	Source	Destination
pg2kafka	postgres_source	kafka_sink

Message example

- Lua script for single message transformation
 - *gofakeit* to generate fake data

```
1 √{
2   "id": 2508,
3   "metadata": {
4     "span_id": "45ead39608ffc918",
5     "trace_id": "ac146df6aa060e043ea390e3756565df"
6   },
7   "payload": {
8     "created_at": "2025-09-03T19:14:55.365Z",
9     "customer_name": "Nickolas Stanton",
10    "event_type": "order_created",
11    "id": "bd009799-000d-4944-b914-9fe49041b55d",
12    "items_count": 9,
13    "quote": "\"Cronut yr tote bag tote bag small batch roof.\""
14  },
15  "created_at": "2025-09-03T19:14:55.365Z"
16 }
```

Replication pitfalls

WAL **bloat** can fill storage,
if replication slot is inactive or not advancing
due to bugs in replication clients or Postgres.

Takeaways

- Avoid microservices if not needed
- Low traffic → *Watermill-SQL, pgx-outbox*
 - High traffic → *Debezium, PeerDB*

Thank you!

PeerDB-io/peerdb

jackc/pglogrepl

nikolayk812/pgx-outbox