

Nikolay Kuznetsov

Senior Software Engineer

Beyond sqlc: teaching AI to generate repositories and integration tests




► About me

- Senior Software Engineer
- Zalando Helsinki 🇫🇮
- C → Java → Kotlin → Go
- Author of *pgx-outbox* and *sqlc++* projects




▶ Goals



- Create repositories on top of *sql/c*-generated queries
- Keep *sql/c*'s strengths: type safety & compile-time checks
- Decouple business logic from *sql/c*-generated records
- Automate repository creation with reference-based AI generation

▶ Repository



- Accepts and returns domain models
- Hides details of SQL, schema, and query libraries
- Maps between DB records and domain models
- Orchestrates transactions across multiple operations

► Cart repository interface

```
type CartRepository interface {  
    // context.Context type is omitted for brevity  
    GetCart(ctx, ownerID string) (domain.Cart, error)  
    AddItem(ctx, ownerID string, item domain.CartItem) error  
    DeleteItem(ctx, ownerID string, productID uuid.UUID) (bool, error)  
}
```

► Domain Cart

```
type Cart struct {
    OwnerID string
    Items   []CartItem
}

type CartItem struct {
    ProductID uuid.UUID // google/uuid
    Price      Money
    CreatedAt time.Time
}

type Money struct {
    Amount    decimal.Decimal // shopspring/decimal
    Currency  currency.Unit     // x/text/currency
}
```

► Cart items table

```
CREATE TABLE IF NOT EXISTS cart_items (  
  owner_id          VARCHAR(255)          NOT NULL,  
  product_id        UUID                  NOT NULL,  
  price_amount       DECIMAL              NOT NULL,  
  price_currency     VARCHAR(3)           NOT NULL,  
  created_at         TIMESTAMP DEFAULT CURRENT_TIMESTAMP NOT NULL,  
  PRIMARY KEY (owner_id, product_id)  
);
```


Working with Postgres in Go

- Pure stdlib database/sql
- squirrel + sqlx
- pgx driver API
- sqlc
- ORMs (~~GORM~~, Ent, Bun): out of scope

Pure stdlib database/sql

- API: *stdlib database/sql*
- Driver: *lib/pq*
- Query builder: none
- Struct mapping: manual scanning

► Pure stdlib database/sql example



```
func (r *repo) GetCart(ctx context.Context, ownerID string) (domain.Cart, error) {  
    // db *sql.DB  
    rows, _ := r.db.QueryContext(ctx,  
        `SELECT product_id, created_at FROM cart_items WHERE owner_id=$1`, ownerID)  
    defer rows.Close()  
  
    // scanning directly into domain model  
    var items []domain.CartItem  
    for rows.Next() {  
        var item domain.CartItem  
        _ = rows.Scan(&item.ProductID, &item.CreatedAt)  
        items = append(items, item)  
    }  
  
    // handle rows.Err()  
  
    return domain.Cart{OwnerID: ownerID, Items: items}, nil  
}
```

Adopting squirrel and sqlx

- jmoiron/sqlx scans to record structs, which need to be mapped to domain models.

- *lib/pq* is maintenance mode and recommends using *jackc/pgx* instead

- API: *stdlib database/sql*
- Driver: *lib/pq*
- Query builder: *Masterminds/squirrel*
- Struct mapping: *jmoiron/sqlx*

► Adopting squirrel and sqlx example

```
func (r *repo) GetCart(ctx context.Context, ownerID string) (domain.Cart, error) {
    query, args, _ := sq.Select("product_id", "created_at").
        From("cart_items").
        Where(sq.Eq{"owner_id": ownerID}).
        PlaceholderFormat(sq.Dollar).ToSql()

    var dbItems []dbCartItem // local struct with `db` tags, aka DB records or models
    // dbx *sqlx.DB
    if err := r.dbx.SelectContext(ctx, &dbItems, query, args...); err != nil {
        return domain.Cart{}, err
    }

    var items []domain.CartItem
    // mapping of dbItems []dbCartItem to domain cart items

    return domain.Cart{OwnerID: ownerID, Items: items}, nil
}
```

Adopting pgx

pgx API advantages:

- Full Postgres feature support (batching, listen/notify, copy)
- Better type handling (JSON, arrays, UUID)
- Advanced connection pool - pgxpool

- API: *jackc/pgx* or *database/sql*
- Driver: *jackc/pgx*
- Query builder: *Masterminds/squirrel*
- Struct mapping: *pgx.CollectRows* method

► Adopting pgx driver and API

```
func (r *repo) GetCart(ctx context.Context, ownerID string) (domain.Cart, error) {
    query, args, _ := sq.Select("product_id", "created_at").
        ... // same as before

    // pool *pgxpool.Pool
    rows, _ := r.pool.Query(ctx, query, args...)

    // collect directly into domain model
    // pgx.CollectRows takes cares of closing rows
    items, _ := pgx.CollectRows(rows, func(row pgx.CollectableRow) (domain.CartItem, error) {
        var item domain.CartItem
        err := row.Scan(&item.ProductID, &item.CreatedAt)
        return item, err
    })

    return domain.Cart{OwnerID: ownerID, Items: items}, nil
}
```


Adopting sqlc

SQL-compiler advantages:

- Compile time safety: schema evolution, queries, type mapping
- Less boilerplate (no rows scanning)
- Separation of SQL and Go code: (queries are in .sql files)

- API: *jackc/pgx*
- Driver: *jackc/pgx*
- Queries: pure SQL
- Struct mapping: automatic to generated sqlc structs

► sqlc in a nutshell



► Generated queries and DB records

```
// Code generated by sqlc. DO NOT EDIT.  
// sqlc v1.29.0
```

```
const GetCart = `-- name: GetCart :many  
SELECT product_id, created_at  
FROM cart_items  
WHERE owner_id = $1  
`
```

```
type GetCartRow struct {  
    ProductID string  
    CreatedAt time.Time  
}
```

```
func (q *Queries) GetCart(ctx context.Context, ownerID string)([]GetCartRow, error) {  
    // generated implementation here  
}
```


► Delegating to sqlc generated queries

```
func (r *repo) GetCart(ctx context.Context, ownerId string) (domain.Cart, error) {  
    // rows []GetCartRow - generated by sqlc  
    // q *db.Queries - generated by sqlc  
    rows, _ := r.q.GetCart(ctx, ownerId)  
  
    // map sqlc rows -> domain items  
    items := mapGetCartRowsToDomain(rows)  
  
    return domain.Cart{OwnerId: ownerId, Items: items}, nil  
}
```

► Generated sqlc structs vs domain models

- sqlc structs \neq domain models
 - They reflect the database schema, not business concepts.
- Avoid using them in business logic
 - Otherwise, business logic would be tightly coupled to the persistence layer.
- Mapping between generated structs and domain models is **needed**

▶ Single Model Couples Your Application



✗ Anti-pattern: The Single Model

Don't give a single model more than one responsibility.

Don't use more than one tag per structure field.

<https://threedots.tech/post/common-anti-patterns-in-go-web-applications/>

▶ Repository methods vs generated sqlc queries

- Repository methods interface align with the **domain**, not SQL
- Repository methods compose multiple queries into a **transaction**
- Repository methods may add **tracing**, metrics, caching, etc

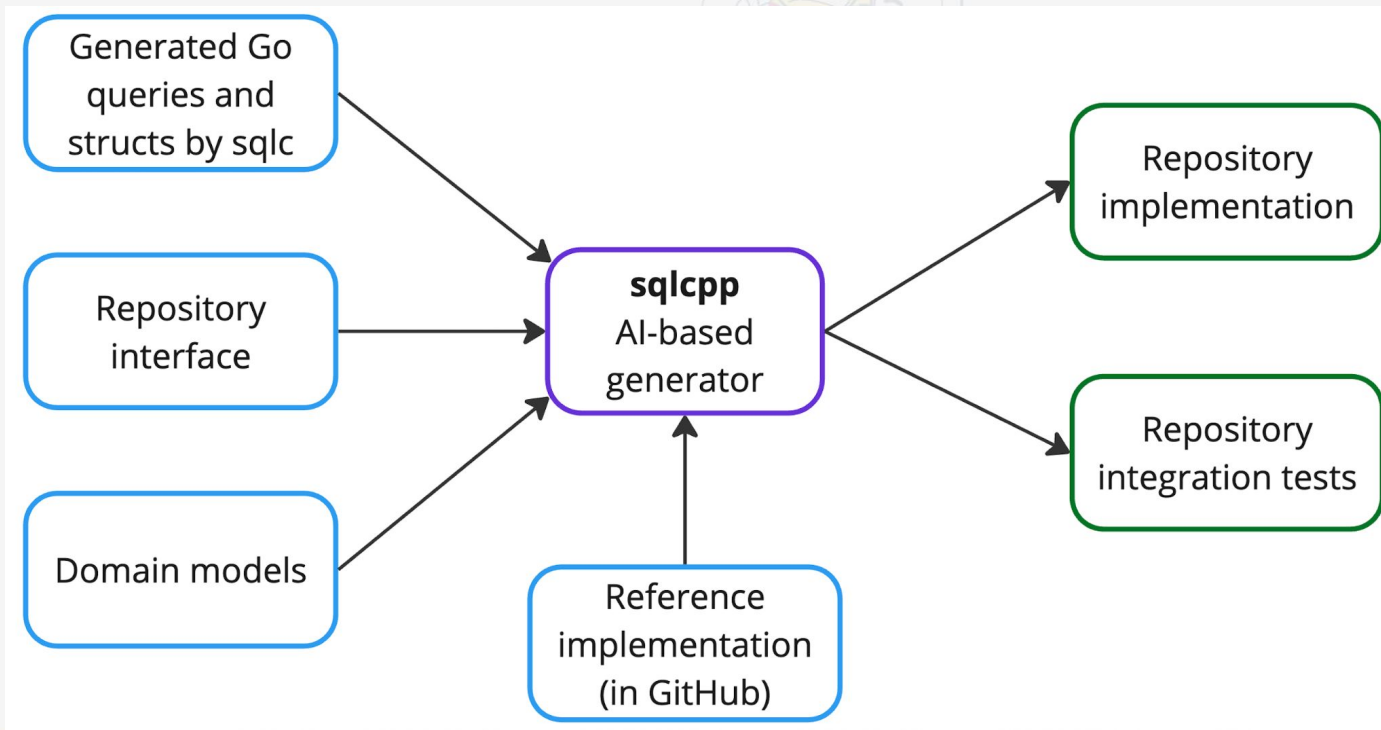
Introducing sqlc++



Let AI generate repository code by:

- Implementing provided (port) interface
- Using domain models and sqlc generated artifacts (queries and structs)
- Leveraging reference-based generation

► In a nutshell



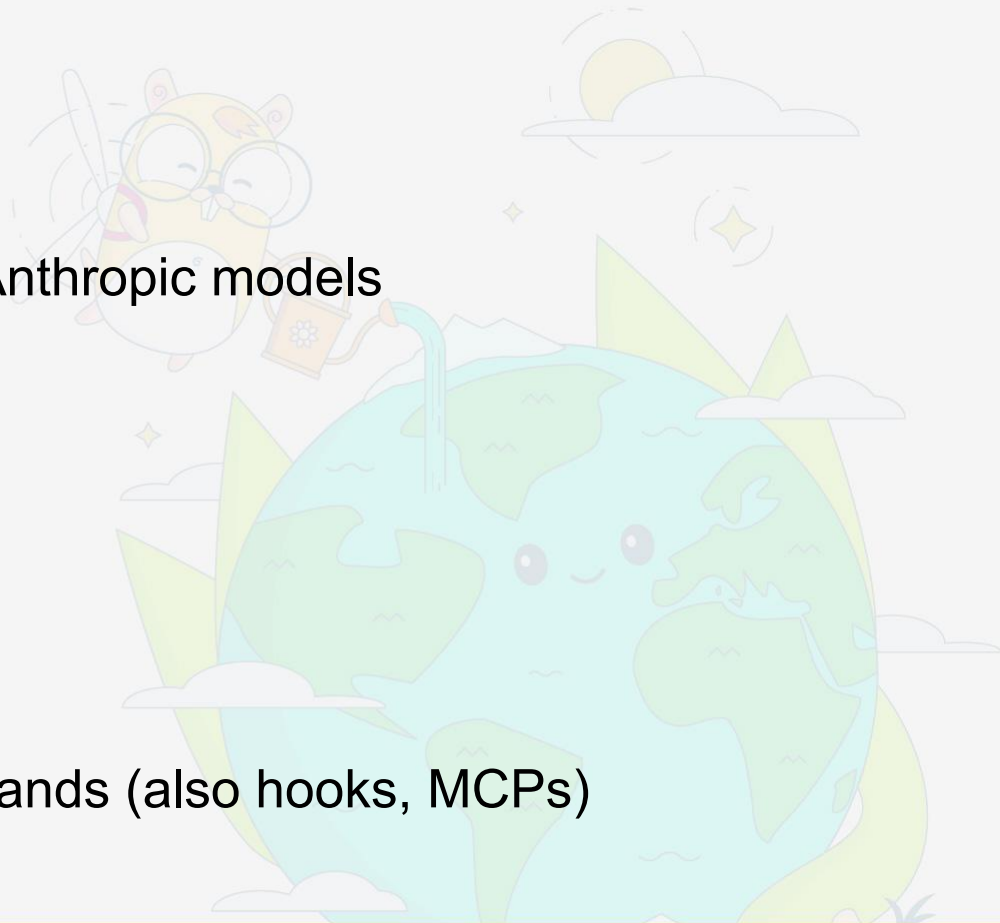
► First version

- Single (huge) request to an LLM
 - Uses *langchain-go* and *text/template*
 - Fast generation (seconds), but inflexible
- 💡 Need to try Claude Code or Crush (ex-OpenCode)



► Claude Code

- Agentic CLI tool powered by Anthropic models
- Specialized for coding tasks
- Can run shell commands
- Can refine results iteratively
- Extensible with custom commands (also hooks, MCPs)



► Current version with Claude Code

- **Custom command** for Claude Code
- Multiple requests to an LLM, more expensive
- Slow generation (minutes), but flexible and powerful
- ***gh*** CLI tool to point to a reference implementation on GitHub

▶ Claude Code's CLAUDE.md

- Similar to *README.md*, can be generated by */init* command
- Claude Code effectively adds it to every request to a LLM
- One can add instructions, i.e. how to fetch a GitHub file content:

```
- `gh api -H "Accept: application/vnd.github.raw"
  repos/nikolayk812/sqlcpp/contents/path/to/file.go`

  to fetch a GitHub file content from
  `https://github.com/nikolayk812/sqlcpp/blob/main/path/to/file.go`
```

▶ Claude Code custom command

- `<project_root>/.claude/commands/generate-repository.md`

Generate a repository file for domain model: `$ARGUMENTS`

Follow these steps:

- Run ``sqlc generate`` command to generate files in ``db`` directory.
- Add generated files in ``db`` directory, ``domain/$ARGUMENTS.go`` and ``port/$ARGUMENTS.go`` to the context.
- Use this file from GitHub as a reference implementation:
``github.com/nikolayk812/sqlcpp/blob/main/repository/order_repository.go``
- Create file ``repository/$ARGUMENTS_repository.go``, it has to satisfy the port interface above.
- Make sure the generated file compiles and tests pass.

▶ Reference implementation project

- Domain model: order
- 2 tables: *orders*, *order_items*
- Various operations: CRUD, search by a filter, soft-delete
- Various field types
 - `uuid.UUID`, `decimal.Decimal`, `currency.Unit`, `time.Time`, `url.URL`
 - `UUID`, `DECIMAL`, `TIMESTAMP`, `JSON`, `JSONB`, `TEXT[]`
- Transactions support

Demo time




Testing repositories



- Testcontainers + Postgres module
- Suite from *stretchr/testify*
- Table tests for each repository method
- Helpful libs: gofakeit, go-cmp

▶ Test Suite

- Suite per a repository, single testcontainer per suite
- Start testcontainer in *SetupSuite* method
- Ryuk terminates unused containers if enabled, otherwise 
- Stop testcontainer in *TearDownSuite* method

► Test Suite

```
type cartRepositorySuite struct {  
    suite.Suite  
  
    repo      port.CartRepository  
    pool      *pgxpool.Pool  
}  
  
func TestCartRepositorySuite(t *testing.T) {  
    suite.Run(t, new(cartRepositorySuite))  
}
```

▶ Setup Suite

```
func (suite *cartRepositorySuite) SetupSuite() {  
    ctx := suite.T().Context()  
  
    connStr, err := startPostgres(ctx)  
    suite.NoError(err)  
  
    suite.pool, err = pgxpool.New(ctx, connStr)  
    suite.NoError(err)  
  
    suite.repo, err = repository.NewCart(suite.pool)  
    suite.NoError(err)  
}
```

▶ Starting Postgres with Testcontainers

```
func startPostgres(ctx context.Context) (string, error) {  
    container, err := postgres.Run(ctx, "postgres:17.6-alpine3.22",  
        postgres.BasicWaitStrategies(),  
        postgres.WithInitScripts("../migrations/01_cart_items.up.sql"),  
    )  
    // handle err  
  
    connStr, err := container.ConnectionString(ctx, "sslmode=disable")  
    // handle err  
  
    return connStr, nil  
}
```

► Random test data with gofakeit

- less boilerplate, edge cases
- possible to pin specific fields for determinism

```
func randomCartItem() domain.CartItem {  
    return domain.CartItem{  
        ProductID: uuid.MustParse(gofakeit.UUID()),  
        Price: domain.Money{  
            Amount: decimal.NewFromFloat(gofakeit.Price(1, 100)),  
            Currency: currency.MustParseISO(gofakeit.CurrencyShort()),  
        },  
    }  
}
```

► Assertions with go-cmp

- readable diff, custom comparers, ignore fields

```
func assertCartItem(t *testing.T, expected, actual domain.CartItem) {
    t.Helper()

    currencyComparer := cmp.Comparer(func(x, y currency.Unit) bool {
        return x.String() == y.String()
    })

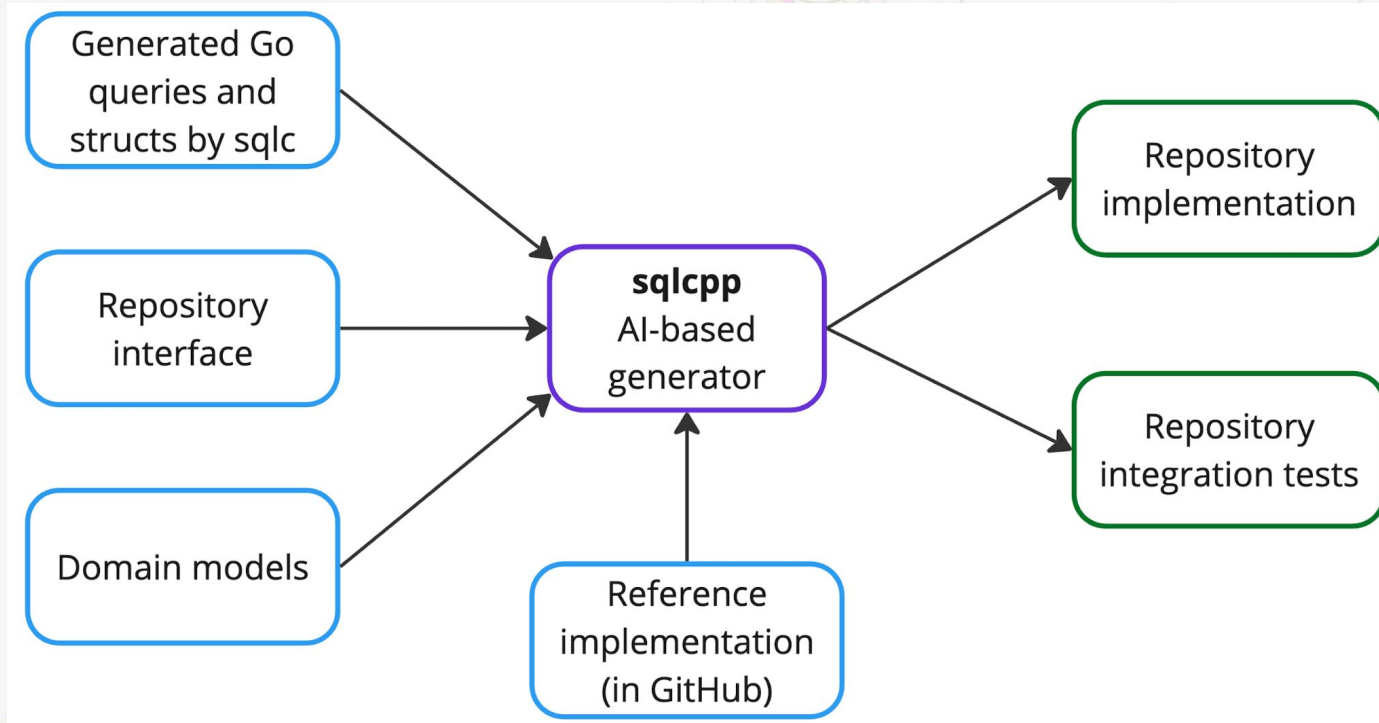
    opts := cmp.Options{
        cmpopts.IgnoreFields(domain.CartItem{}, "CreatedAt"),
        currencyComparer,
    }

    diff := cmp.Diff(expected, actual, opts)
    assert.Empty(t, diff)
}
```


Demo results



► sqlc++ architecture recap



► Takeaways



- Adopt sqlc, avoid use of generate structs in business logic
- Adopt agentic CLI tool: (Claude Code, Crush + GPT-OSS in Ollama)
- Reference-based generation provides more stable results
- Repository boilerplate is automated
- Schema, queries, and domain models remain manual

Thank you!

Q&A



Nikolay Kuznetsov



nikolayk812



nkuznetsov



github.com/nikolayk812/sqlcpp

